

Introduction to  
Automatic Adjoint Differentiation  
- AAD -  
in Machine Learning  
and Finance

Antoine Savine

<http://antoinesavine.com>

<http://github.com/asavine/CompFinance/Intro2AADinMachineLearningAndFinance.pdf>

<http://tinyurl.com/savineAadTalk>

follow asavine on GitHub and watch the repo to be notified of updates

# Introduction

# Adjoint Differentiation - AD -

- Algorithm to compute **all** differentials of a scalar function **quickly** and **accurately**

- **Quickly** (*constant time*):  
Compute **all** differentials in a time similar to **one** function evaluation
- **Accurately**:  
**Analytic** differentiation,  
accurate to machine precision

$$f: \mathbb{R}^D \rightarrow \mathbb{R}$$
$$x = \begin{pmatrix} x_1 \\ \dots \\ x_D \end{pmatrix} \rightarrow y = f(x)$$

$$AD: compute: \frac{\partial f}{\partial x} = \left( \frac{\partial f}{\partial x_1}, \dots, \frac{\partial f}{\partial x_D} \right) \text{ in } O(1)$$

- Application:  
whenever we need to compute **many** derivatives of **one** result
- AD is *only* about speed yet a massive difference in Machine Learning and Finance – without AD:
  - Deep ANNs could not learn their parameters in reasonable time
  - Financial risks could not be computed in real time

# Application: model fitting

- Model:  $\hat{y} = f(x; \mathcal{G})$ 
  - Makes a prediction  $\hat{y}$  from a set of inputs  $x = (x_1, \dots, x_n)^T$
  - Given parameters  $\mathcal{G} = (\mathcal{G}_1, \dots, \mathcal{G}_D)^T$
- Learns its parameters by minimizing a *cost* function
  - E.g. sum of squared errors:  $C(\mathcal{G}) = \sum_{i=1}^m \left[ f(x^{(i)}; \mathcal{G}) - y^{(i)} \right]^2 = \sum_{i=1}^m c_i(\mathcal{G})$  (  $c_i$  : error or *loss* on training example  $i$  )
  - Over a training set of  $m$  labelled examples:  $(x^{(i)}, y^{(i)})_{1 \leq i \leq m}$  (  $x^{(i)} \in \mathbb{R}^n, y^{(i)} \in \mathbb{R}$  )
- Must compute gradient of the cost function for optimization algorithm:  $\frac{\partial C}{\partial \mathcal{G}} = \left( \frac{\partial C}{\partial \mathcal{G}_1}, \dots, \frac{\partial C}{\partial \mathcal{G}_D} \right)$ 
  - Standard case (sum of squared errors):  $\frac{\partial C}{\partial \mathcal{G}} = \sum_{i=1}^m \frac{\partial c_i}{\partial \mathcal{G}}$  ,  $\frac{\partial c_i}{\partial \mathcal{G}} = 2c_i \frac{\partial f(x_i; \mathcal{G})}{\partial \mathcal{G}}$

# Training a deep learning model

- Multi-Layer Perceptron (MLP)  
simplest and most common deep learning model:


- Prediction: feed-forward equations  $a^{[l]} = g^{[l]}(W^{[l]}a^{[l-1]} + b^{[l]})$
- Training set:  $m$  inputs  $x^{(i)}$ , each a vector in dimension  $n_0$  with corresponding (scalar) labels  $y^{(i)}$

- Train the network: learn (fit) parameters  $W^{[1]}, b^{[1]}, \dots, W^{[L]}, b^{[L]}$  to minimise cost = sum of errors (losses) on the training set  $c^{(i)} = (\hat{y}^{(i)} - y^{(i)})^2$
- Then apply trained model to make predictions on new examples with unknown label

- Requires an iterative minimization algorithm

- On each iteration, we must provide  $C$ :  $m$  predictions through a deep network

- And all their derivatives to parameters  $\frac{\partial c^{(i)}}{\partial W_{j,k}^{[l]}}, \frac{\partial c^{(i)}}{\partial b_j^{[l]}}$



# Calibrating a financial pricing model

- Calibration of a pricing model to market data:
  - Prediction: valuation of a financial product  $\hat{y} = \text{product value} = f(\text{product parameters } x; \text{model parameters } \theta)$
  - Training set:  $m$  products, each with  $n$  known features  $x^{(i)}$  and market price  $y^{(i)}$
  - Calibration: fit parameters  $\theta = (\theta_1, \dots, \theta_D)^T$  (e.g. volatilities) to minimise cost = sum of errors (losses) on the training set
  - Then apply calibrated model to price other products which prices are unknown
- Minimization algorithms (gradient descent, conjugate gradients, pseudo-Newton, Levenberg-Marquardt...)
  - Requires derivatives  $\frac{\partial C}{\partial \theta} = \left( \frac{\partial C}{\partial \theta_1}, \dots, \frac{\partial C}{\partial \theta_D} \right)$  on each iteration, we have as many differentials as we have parameters
  - in addition  $C$  may be expensive to evaluate, must value all  $n$  products in the calibration set

# Notes on calibration

- Calibration is considered best practice in financial derivatives (as opposed to estimation of parameters)
- To understand exactly why, see  
<https://www.slideshare.net/AntoineSavine/60-years-birthday-30-years-of-ground-breaking-innovation-a-tribute-to-bruno-dupire-by-antoine-savine>
- Calibration is only similar to training ML models in appearance
  - ML models are *entirely* trained on data and make predictions on new examples *drawn from the same distribution*
  - “Interpolation” problem
  - Financial models have a strong structure: arbitrage-free, realistic dynamics, some parameters (e.g. correlation) typically estimated
  - Often calibrated on European options to value more complicated *exotic* options
  - “Extrapolation” problem



# Application: model risk

- Financial model:  $value = f(\text{product features } x; \text{model parameters } \mathcal{G})$
- Model parameters:
  - Today's underlying asset prices
  - Volatilities, correlations
  - Mean-reversions
  - Etc.
- Model risks:  $\frac{\partial f}{\partial \mathcal{G}}$  with potentially massive number of parameters
  - Black-Scholes: spot, volatility, rate
  - Dupire: two-dimensional local volatility surface, with stochastic volatility: vol of vol, correlation spot/vol
  - Interest Rate Models: initial yield curve + spread curve(s) + volatility curve or surface (+ mean-reversion...)
  - Multi-currency models (Baskets, CVA...): all of the above for all currencies + forex rates and their volatilities + giant correlation matrix
  - Typically in the thousands or tens of thousands
  - Valuation generally expensive: Monte-Carlo simulations

# Market risk

- Usually, we are not interested in risks to model parameters but risks to market variables

E.g. Dupire: not local volatilities but implied volatilities

Model parameters are generally calibrated to market variables

$a$ : vector of  $k$  market variables (underlying asset prices, rates, spreads, implied volatilities...)

$b$ : vector of  $n$  model parameters (calibrated to market variables so  $b = c(a), c: \mathbb{R}^k \rightarrow \mathbb{R}^n$ )

$g$ : valuation function of model (generally numerical such as Monte-Carlo simulation)  $V = g(b), g: \mathbb{R}^n \rightarrow \mathbb{R}$

$f$ : valuation function of market  $V = g(b) = g[c(a)] = f(a), f: \mathbb{R}^k \rightarrow \mathbb{R}$

- Risk report computes model risks  $\frac{\partial g}{\partial b}$  and/or market risks  $\frac{\partial f}{\partial a}$

- Market risks computed from model risks:  $a \in \mathbb{R}^k \xrightarrow{c} b \in \mathbb{R}^n \xrightarrow{g} v \in \mathbb{R}$  so by the chain rule  $\frac{\partial f}{\partial a} = \frac{\partial g}{\partial b} \frac{\partial c}{\partial a}$

- To compute  $\frac{\partial c}{\partial a}$  and market risks  $\frac{\partial f}{\partial a} = \frac{\partial g}{\partial b} \frac{\partial c}{\partial a}$  knowing model risks  $\frac{\partial g}{\partial b}$ , see for instance:

[http://papers.ssrn.com/sol3/papers.cfm?abstract\\_id=3262571](http://papers.ssrn.com/sol3/papers.cfm?abstract_id=3262571)

- In this presentation, we focus on model risks  $\frac{\partial g}{\partial b}$

# Differentiation

- In all these applications (and many others)
  - We have a scalar function  $f$  of many inputs  $f: \mathbb{R}^D \rightarrow \mathbb{R}$
  - This function is typically expensive to evaluate  
Losses of deep neural nets on large training set or a Monte-Carlo valuation may take several seconds, even on parallel hardware
  - It typically takes a large number  $D$  of inputs  
Thousands of weights for deep neural nets, or thousands of market/model variables for financial valuation
  - We must compute its sensitivities to all its inputs  $\frac{\partial f}{\partial x} = \left( \frac{\partial f}{\partial x_1}, \frac{\partial f}{\partial x_2}, \dots, \frac{\partial f}{\partial x_D} \right)$
- Conventional differentiation is of linear complexity in the number of inputs/differentials
  - To compute **one** differential takes time similar to **one** evaluation of the function
  - For example, with (one-sided) finite differences  $\frac{\partial f}{\partial x_i} \approx \frac{f(x + \varepsilon e_i) - f(x)}{\varepsilon}, e_i = (\delta_{ij})_{1 \leq j \leq n}$
  - We have  $(D+1)$  function evaluations to compute  $D$  differentials
  - Illustration: if one function evaluation takes 1sec and we have 1,000 inputs  
It takes 1,001sec (15+ minutes) to evaluate its gradient once
  - This is not viable, a different technology is necessary

# Adjoint Differentiation - AD -

- Algorithm to compute analytically all the differentials of  $f$ :  $\frac{\partial f}{\partial x} = \left( \frac{\partial f}{\partial x_1}, \frac{\partial f}{\partial x_2}, \dots, \frac{\partial f}{\partial x_D} \right)$ 
  - In a time independent of  $D$
  - With an efficient implementation, it takes 4 to 10 times **one** evaluation of  $f$  to compute **all** its differentials
  - So AD can compute the 1,000 differentials of a function that takes 1sec in ~4 to 10sec as opposed to 15min+
  - Without loss of accuracy (if anything, AD differentials are *more* accurate)
- Hence the importance of this technology today
  - In finance, AD computes thousands of complex risks accurately and very quickly
  - In machine learning, AD computes the gradient of cost functions to many parameters in constant time
  - The same technology powers banks to estimate financial risks and deep nets to learn their weights
  - AD
    - Allows Banks to compute massive amount of regulatory risks in real time
    - Powers your telephones to learn to recognize you and your friends in reasonable time
    - Is largely credited for the recent and spectacular successes in the field of deep learning

# A brief history of AD

- Invented in the 1960s (Wengert, 1964)
- Called “holy grail” of sensitivity computation (Griewank, 2012)
- Classified as on the 30 greatest numerical algorithms of the 20<sup>th</sup> century (Trefethen, 2015)
- Did not take firm hold in the computer science community until 1980
- Applied in deep learning to efficiently compute gradients of cost functions in learning algorithms
- Known under the names  
Back-Propagation (or simply back-prop, in deep learning),  
reverse differentiation, backward differentiation, adjoint accumulation, algorithmic differentiation, etc.
- Not adopted in Finance until 2006 (Giles and Glasserman’s “Smoking Adjoint”)
- Large scale implementation in Danske Bank with parallel Monte-Carlo simulations won In-House System of the Year 2015 Risk Award
- Widely adopted today for risk and calibration
- Delays in adoption mainly due to complexity, lack of teaching material and challenges in practical implementation (Andersen, 2018)


# Automatic Adjoint Differentiation - AAD -

- Automatic implementation of AD
  - Developers only produce the evaluation code for the function  $f$
  - The framework automatically produces (constant time) differentiation “code” for  $\partial f / \partial x$
  - So developers don’t need to write (complicated and error prone) AD code
  - And differentiation code is automatically updated when evaluation code changes
- Many commercial and open source AAD frameworks exist
  - Some are generic, some specialize in finance or machine learning
  - Some are free, some as extremely expensive
  - Work in a variety of programming languages: Python, C++, ...
  - Work either with source transformation, operator overloading or by letting clients build evaluation graphs
- Example: TensorFlow, one of most popular framework in machine learning
  - Written in C++/CUDA (GPU)
  - With APIs in Python, JavaScript, Java, Go, Swift...
  - Allow clients to create evaluation graphs
  - TensorFlow evaluates the graphs and automatically applies AD to compute derivatives through the graphs

# AAD: challenges

- However it is challenging to develop or even use AAD frameworks efficiently:
  - AD is somewhat “mind twisting” and generally poorly understood
  - Memory management: AAD consumes vast amounts of RAM
  - RAM, cache and parallel efficiency are hard to achieve
  - Aspects of practical implementation are topics of advanced research
- Blindly applying a framework almost never works
  - Not quite as easy as coding an evaluation and sending it to the framework for differentiation
  - Requires a deep understanding of the theory and implementation
  - Otherwise often results in slow differentiation, incorrect results, crashes, etc.
- This presentation covers the bases of AD and AAD and introduces the book:  
**Modern Computational Finance: AAD and Parallel Simulations**  
which explains in deep detail all aspects of AAD and its practical implementation


# AAD book



- Written by (some of) the people who wrote Danske Bank's award winning systems
- Prefaced by Leif Andersen, read preface on [researchgate.net/publication/328042479\\_Modern\\_Computational\\_Finance\\_AAD\\_and\\_Parallel\\_Simulations](https://researchgate.net/publication/328042479_Modern_Computational_Finance_AAD_and_Parallel_Simulations)
- Teaches:
  - The modern design and efficient implementation of financial simulation libraries
  - The implementation of parallel simulation libraries
  - And of course AAD
- AAD is covered in deep detail:
  - Conceptual and mathematical foundations
  - Practical implementation
  - Memory management and check-pointing
  - Application in the context of large Monte-Carlo simulations, including in parallel
  - Application to model and market risks
  - Cutting-edge implementation with meta-programming and expression templates
  - And much more
- Ships with complete, professional code in C++
  - Explained in deep detail in the book
  - Freely available on the GitHub repo <http://www.github.com/asavine/CompFinance/wiki>



# AAD book reviews



From Amazon: <https://www.amazon.com/gp/product/1119539455>

It would not be much of an exaggeration to say that Antoine Savine's book ranks as the 21st century peer to Merton's 'Continuous-Time Finance'

**Vladimir Piterbarg**

This book [...] addresses the challenges of AAD head on. [...] The exposition is [...] ideal for a Finance audience. The conceptual, mathematical, and computational ideas behind AAD are patiently developed in a step-by-step manner, where the many brain-twisting aspects of AAD are demystified. For real-life application projects, the book is loaded with modern C++ code and battle-tested advice on how to get AAD to run for real. [...] Start reading!

**Leif Andersen**

An indispensable resource for any quant. Written by experts in the field and filled with practical examples and industry insights that are hard to find elsewhere, the book sets a new standard for computational finance.

**Paul Glasserman**

# Overview

## AAD: Demonstration

1. Adjoint Differentiation for Deep Learning
  1. A brief introduction to Artificial Neural Networks (ANN)
  2. Back-Propagation through ANNs
2. Adjoint Differentiation for arbitrary calculations
  1. Evaluation graphs
  2. Back-Propagation through evaluation graphs
3. Automatic Adjoint Differentiation in (simple) code
  1. AAD with operator overloading
  2. AAD over Monte-Carlo simulations

# Demonstration: Dupire's model (1992)

- Extended Black & Scholes dynamics (in the absence of rates, dividends etc.):  $\frac{dS}{S} = \sigma(S, t) dW$
- Calibrated with Dupire's celebrated formula:  $\sigma(K, T) = \frac{2 \frac{\partial C}{\partial T}}{\frac{\partial^2 C}{\partial K^2}}$  with  $C(K, T)$  = call prices of strike  $K$ , maturity  $T$
- Implemented with a (bi-linearly interpolated) local volatility matrix:  $\sigma_{ij} = \sigma(S_i, T_j)$
- Volatility matrix: 30 spots (every 5 points 50 to 150) and 36 times (every month from 0 to 3y)  
we have 1,080 volatilities + 1 initial spot (=100) = 1,081 model parameters
- Valuation of a 3y (weekly monitored) barrier option strike  $K=120$ , barrier  $B=150$ :  $v(S_t, t) = E \left[ (S_T - K)^+ 1_{\{\max(S_{T_1}, \dots, S_{T_K}) < B\}} \middle| S_t \right]$
- Solved with Monte-Carlo or FDM over the equivalent PDE (from Feynman-Kac's theorem)
- We focus on Monte-Carlo simulations here: 500,000 paths, 156 (weekly) time steps

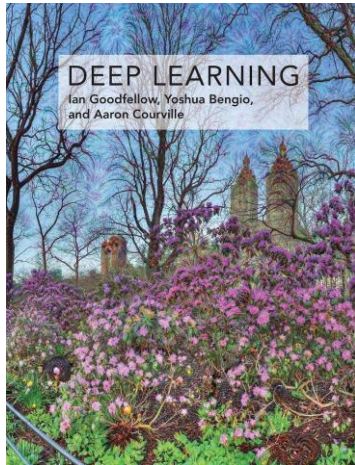
# Demonstration: Results

- Implementation
  - C++ code exported to Excel (see tutorial on <http://www.github.com/asavine/CompFinance/wiki>)
  - Generic library design with efficient implementation (Chapter 6)
  - Parallel implementation (Chapters 3 and 7)
  - Sobol quasi-random numbers (Chapters 5 and 6)
  - Advanced AAD with expression templates (Chapter 15)
- Hardware: quad-core laptop (surface book 2, 2017)
- Performance:
  - One evaluation with 500,000 paths over 156 time steps take ~0.8sec
  - We have 1,081 risk sensitivities, take about 15 minutes to produces model risk report with linear differentiation
  - With AAD the 1,081 differentials are produced in ~1.5sec

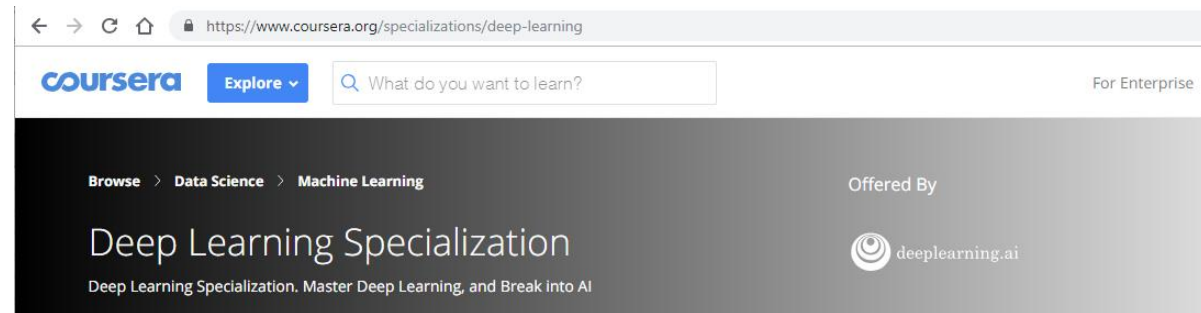
# Deep Learning

# Neural networks and deep learning

- Adjoint differentiation is best known and explained in the context of deep learning
- We therefore *briefly* introduce deep learning
- (Much) deeper presentations are found in:



Goodfellow's book




Andrew Ng's videos on Coursera

...and many other books and resources

# Linear regression: prediction

- Linear model (joint Gaussian assumptions):  $\hat{y} = E[y|x] = b + \sum_{i=1}^n w_i x_i = b + w^T x$

- Parameters:  $b \in \mathbb{R}$  and  $w \in \mathbb{R}^n$




# Linear regression and classification (1)

- In *regression* problems we predict real numbers

Example: predict return from volatility

- Each point is a stock in the S&P500
- Horizontal axis is annual volatility
- Vertical axis is annual return
- Estimated with daily data, 2013-2018
- Evidence of (small) risk premium
- Poor regression quality





# Linear regression and classification (2)

- Alternatively, *classification* problems predict discrete categories

Example: identify animals in pictures: 0: not an animal, 1: cat, 2: dog, 3: bird, 4: other animal

picture 1920x1080x24bit colour



vector of 1920x1080 pixels

$$\in [0, 2^{24} - 1]$$

$$x = \begin{pmatrix} x_1 \\ \dots \\ x_{2,073,600} \end{pmatrix}$$

softmax regression

$$\hat{y} = s(b + wx) = \begin{bmatrix} \Pr(\text{no animal}) \\ \Pr(\text{cat}) \\ \Pr(\text{dog}) \\ \Pr(\text{bird}) \\ \Pr(\text{other animal}) \end{bmatrix}$$

Parameters:

b: vector of dimension 5

w: matrix 5 x 2,073,600

s: softmax function

$s: \mathbb{R}^5 \rightarrow (0,1)^5$  summing to 1

$$s(z) = \begin{bmatrix} \frac{e^{z_i}}{\sum e^{z_j}} \end{bmatrix}$$

- In this presentation, we stick with regression
- Everything that follows generalises easily to classification and the equations remain essentially identical
- See literature for details, for example Stanford's CS229 on <http://cs229.stanford.edu/syllabus.html>

# Linear regression: training

- Training set of  $m$  examples  $x^{(1)}, \dots, x^{(m)}$  (each a vector in dimension  $n$ ) with corresponding labels  $y^{(1)}, \dots, y^{(m)} \in \mathbb{R}$

- Learn parameters:  $b \in \mathbb{R}$  and  $w \in \mathbb{R}^n$  by minimizing prediction errors (cost function)  $C(b, w) = \sum_{i=1}^m \left( \underbrace{b + w^T x^{(i)}}_{=\hat{y}^{(i)}} - y^{(i)} \right)^2$

- Note that this is the same as maximizing (log) likelihood under Gaussian assumptions, hence:

$b^*, w^* = \arg \min C(b, w)$  are the maximum likelihood estimators (MLE) of the parameters

- $b^*$  and  $w^*$  are found analytically, solving for  $\frac{\partial C}{\partial b} = 0$  and  $\frac{\partial C}{\partial w} = 0$

- Result (“normal equation”):  $\begin{pmatrix} b \\ w_1 \\ w_2 \\ \dots \\ w_n \end{pmatrix} = (X^T X)^{-1} X^T Y$  where  $X = \begin{pmatrix} 1 & x_1^{(1)} & \dots & x_n^{(1)} \\ 1 & x_1^{(2)} & \dots & x_n^{(2)} \\ \dots & x_1^{(i)} & x_j^{(i)} & x_n^{(i)} \\ 1 & x_1^{(m)} & \dots & x_n^{(m)} \end{pmatrix}$  and  $Y = \begin{pmatrix} y^{(1)} \\ y^{(2)} \\ \dots \\ y^{(m)} \end{pmatrix}$


# Lin reg only captures linear functions

- Example:

- Predict the future price in 1y of a European call strike 100, maturity 2y
- By regression of the payoff in 2y:  $y^{(i)} = (S_{2y}^{(i)} - K)^+$
- Over the underlying asset price in 1y:  $x^{(i)} = S_{1y}^{(i)}$
- With a training set of  $m$  paths  $(S_{1y}^{(i)}, S_{2y}^{(i)})_{1 \leq i \leq m}$

generated under Black & Scholes' model (spot = 100, volatility of 20%)


- We know the exact solution is given by Black and Scholes' formula so we can assess the quality of the regression
- Linear regression obviously fails to approximate the correct function because it cannot capture non-linearities



# Basis function regression

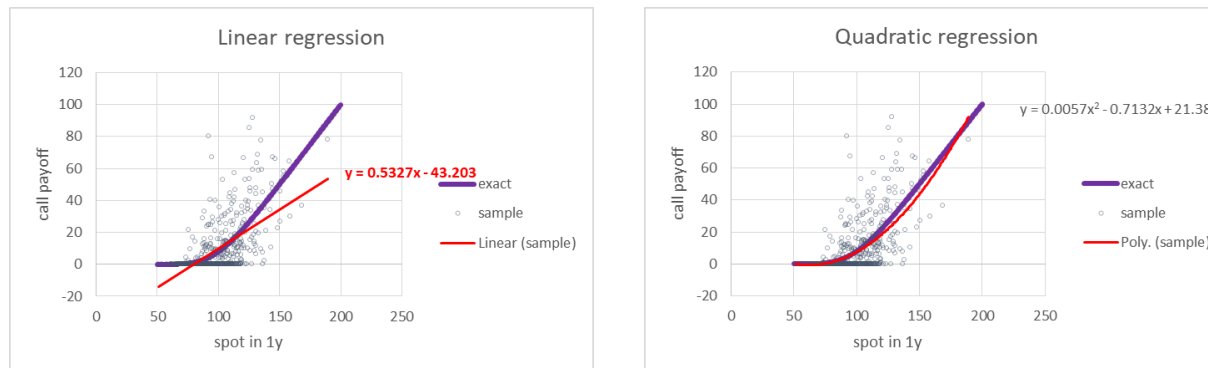
- Solution: regress not on  $x$  but on basis functions  $a$  of  $x$
- Example: polynomial regression
  - Basis functions = monomials ( $x$  is in dimension  $n_0$ ):
    - 1<sup>st</sup> degree: all the xs  $a_i = x_i$
    - 2<sup>nd</sup> degree: all the squares  $a_i = x_j^2$  and pair-wise products  $a_i = x_i x_j$
    - 3<sup>rd</sup> degree: all the cubes  $a_i = x_j^3$  and pair-wise  $a_i = x_j x_k^2$  and triplet-wise  $a_i = x_j x_k x_l$
    - Etc.
- Prediction in two steps:
  - Start with the vector  $x$  of  $n_0$  features
  - Compute vector of  $n_1$  basis functions:  $a = \varphi(x)$
  - Predict linearly **in the basis functions**:  $\hat{y} = b + w^T a$
- Training: identical to linear regression on  $a$  in place of  $x$

$$\begin{pmatrix} b \\ w_1 \\ w_2 \\ \dots \\ w_{n_1} \end{pmatrix} = (A^T A)^{-1} A^T Y \quad \text{where} \quad A = \begin{pmatrix} 1 & a_1^{(1)} & \dots & a_{n_1}^{(1)} \\ 1 & a_1^{(2)} & \dots & a_{n_1}^{(2)} \\ \dots & a_1^{(i)} & a_j^{(i)} & a_{n_1}^{(i)} \\ 1 & a_1^{(m)} & \dots & a_{n_1}^{(m)} \end{pmatrix} \quad \text{and} \quad a^{(i)} = \varphi(x^{(i)})$$



# Basis function regression: performance

- Works nicely: quadratic regression of  $(S_{T_2} - K)^+$  over  $S_{T_1}$  and  $S_{T_2}$  approximates Black & Scholes' formula well in simulated example



- Remarkable how the algorithm manages to detect Black & Scholes' pattern in noisy data
- Mathematically:
  - Combinations of polynomials can approximate any smooth function to arbitrary precision
  - Hence, detect any (smooth) non-linear pattern in data
  - *But only with a large number of basis functions / high polynomial degree*

# Curse of dimensionality

- How many monomials in a  $p$ -degree polynomial regression?

Number  $n_1$  of basis functions (dimension of  $a$ ) grows exponentially in number  $n_0$  of features (dimension of  $x$ )

Precisely:  $n_1 = \frac{(n_0 + p)!}{n_0! p!} - 1$

- Quadratic regression:  $n_1 = \frac{(n_0 + 2)(n_0 + 1)}{2} - 1$  dimension grows e.g. from 10 to 65, from 100 to 5,151, from 1,000 to 501,500
- Cubic regression:  $n_1 = \frac{(n_0 + 3)(n_0 + 2)(n_0 + 1)}{6} - 1$  dim grows 10 to 286, 100 to 176,851, 1,000 to 167,668,501
- In general, number of basis functions increases exponentially in dimension
- “Rule of ten”:  
to avoid overfitting small data set with many parameters, we need an amount of data (training examples) of approximately  $m > 10n$
- Hence, basis function regression requires amount of data exponential in (original) dimension  $n_0$
- Basis function regression works nicely in low dimension but doesn’t scale (think of image processing where  $n_0 = 1920 \times 1080 > 2M$ )


# Overfitting and the rule of ten

- If you fit a linear model with  $n$  free parameters to  $n$  data points
  - You will find a perfect fit
  - But the model may not generalize well to new examples
  - Because it captured the noise of the particular training set, “overfitting”
- For error analysis (“variance-bias trade-off”), see:  
<https://www.quora.com/Does-neural-network-generalize-better-if-it-is-trained-on-a-larger-training-data-set-Are-there-any-scientific-papers-showing-it/answer/Antoine-Savine-1>
- Solutions exist such as regularization  
See **Stanford’s Machine Learning class on Coursera** or **Bishop’s Pattern Recognition and Machine Learning**
- But the most effective means to avoid overfitting is increase the size of the learning set
- An empirical rule of thumb called “rule of ten”  
recommends a data set of at least 10x the number of fitted parameters

# Idea (ANN): learn basis functions from data

## BASIS FUNCTION REGRESSION

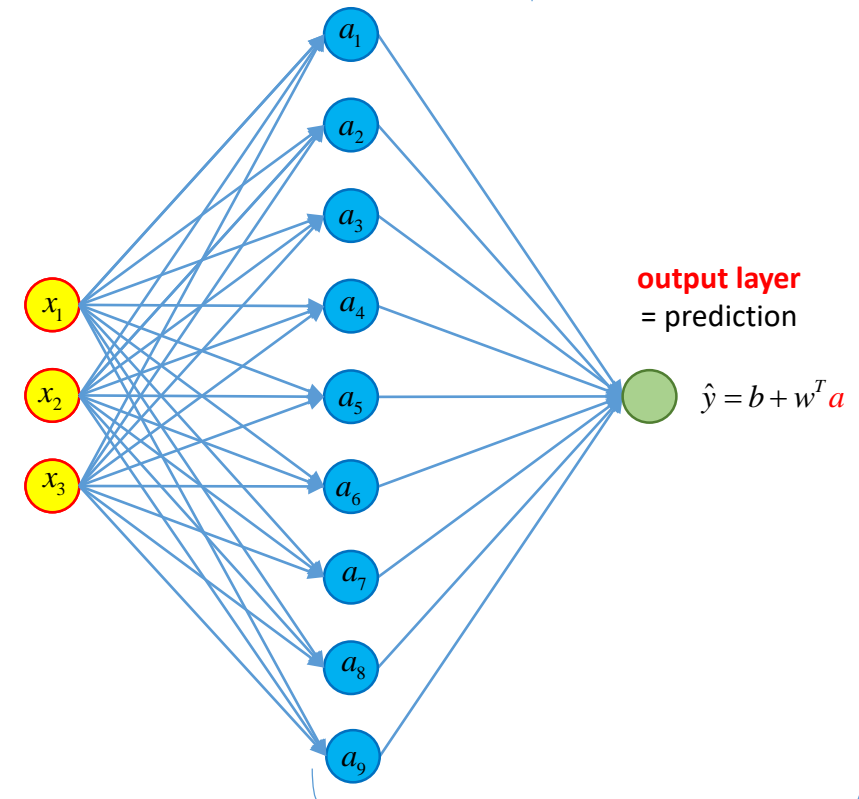
hidden layer = **arbitrary** set of basis functions  
pre-processing = not part of the model, no parameters



linear regression over basis functions

## ANN

hidden layer = **learned** set of basis functions  
part of the model, subject to trainable parameters



linear regression over basis functions  
**identical**



# Shallow ANN: parametric basis functions

- How do we learn basis functions  $a$  from the data?
  - Parametric family of basis functions:  $a_i = \varphi(x; \theta_i) \rightarrow$  each basis function is the *same function of  $x$* , with *different parameters  $\theta$*
  - Learn parameters  $\theta$  (along with the weights  $w$  and bias  $b$  of regression) by minimization of the cost (e.g. sum of squared errors)
- Most common choice of parametric basis functions: “perceptron”
  - Each basis function is a non-linear scalar function  $g$  of a linear combination of  $x$ s (different for each basis function):  $a_i = g(w_i^T x + b_i)$
  - Where each  $w_i, 1 \leq i \leq n_1$  is a vector in dimension  $n_0$  and each  $b_i, 1 \leq i \leq n_1$  is a real number
  - Hence,  $W = \begin{pmatrix} w_1^T \\ \dots \\ w_{n_1}^T \end{pmatrix}$  is a  $n_1 \times n_0$  matrix and  $b = \begin{pmatrix} b_1 \\ \dots \\ b_{n_1} \end{pmatrix}$  is a  $n_1$  dimensional vector
  - Then the vector of basis functions is  $a = g(Wx + b)$  where  $g$  is a scalar function applied element-wise to the  $n_1$  vector  $z = Wx + b$

# ANN: prediction

- Prediction in two steps

- Step 1: hidden layer (basis functions) **inputs  $x \rightarrow$  basis functions  $a$**   $a = g(W^{[1]}x + b^{[1]})$ 
  - Parameters:  $W^{[1]}$  matrix in dimension  $n_1 \times n_0$  and  $b^{[1]}$  vector in dimension  $n_1$
  - The *activation* function  $g$  is scalar and applied element-wise on the  $n_1$  vector  $z^{[1]} = W^{[1]}x + b^{[1]}$
  - $g$  must be non-linear or we are back to linear regression  
(linear regression on linear functions is identical to linear regression on inputs)

- Step 2: regression **basis functions  $a \rightarrow$  prediction**  $\hat{y} = w^{[2]}a + b^{[2]}$ 
  - Parameters:  $w^{[2]}$  column vector in dimension  $n_1$  and  $b^{[2]}$  real number
  - $\hat{y}$  is a real number

- Unified notation for both steps = **feed-forward equation**  $a^{[l]} = g^{[l]}(W^{[l]}a^{[l-1]} + b^{[l]})$

- With the notations
 

$a^{[0]} = x$	input layer	dim $n_0$	
$a^{[1]} = a$	hidden layer	dim $n_1$	
$\hat{y} = a^{[2]}$	output layer	dim $n_2 = 1$	
$g^{[1]} = g, g^{[2]} = id$	hidden and output activations		




 and parameters  $W^{[l]}$  of dim  $n_l \times n_{l-1}$  and  $b^{[l]}$  of dim  $n_l$  for  $l = 1, 2$

# Choice of activation function

- The hidden activation function  $g^{[1]}$  must be non-linear – the most common choices are:

- “Cheap Gaussians”
  - sigmoid  $\sigma(x) = \frac{1}{1 + \exp(-x)}$
  - tanh  $t(x) = \frac{\exp(x) - \exp(-x)}{\exp(x) + \exp(-x)}$

- Derivatives and integrals of sigmoid
  - $\sigma'(x) = \sigma(x)(1 - \sigma(x))$
  - $\int_{-\infty}^x \sigma(x) = \log(1 + \exp(x))$  (“softPlus”)
- Rectified Linear Units (ReLU)**  $r(x) = (x)^+$




- The output activation function  $g^{[2]}$  is:
  - Identity  $g^{[2]}(x) = x$  for regression
  - SoftMax for classification

# ANN: computation graph

Input layer  $l=0$       Hidden layer  $l=1$       Output layer  $l=2$

$$a^{[0]} = x \quad \longrightarrow \quad a^{[1]} = g^{[1]} \left( W^{[1]} a^{[0]} + b^{[1]} \right) \longrightarrow a^{[2]} = g^{[2]} \left( W^{[2]} a^{[1]} + b^{[2]} \right) = \hat{y}$$

$n_0$  units       $n_1$  units       $n_2=1$  units



feed-forward equation

$$a^{[l]} = g^{[l]} \left( W^{[l]} a^{[l-1]} + b^{[l]} \right)$$

## Computation complexity (time)

To the leading order (ignoring activations and additions)

We have 2 matrix by vector products  $W^{[1]} a_0$  and  $W^{[2]} a_1$

$n_1 \times n_0 \quad n_0 \qquad n_2 \times n_1 \quad n_1$

Hence quadratic complexity  $n_1 n_0 + n_2 n_1$

To simplify when all layers have  $n$  units, complexity  $\sim 2n^2$

More generally, with  $L$  hidden layers, complexity  $\sim Ln^2$

Examples:

- With 1,000 features/basis functions and 2 layers  
complexity  $\sim 2M$ , fraction of a second on modern CPU
- With 1M features/basis functions and 100 layers  
(not unrealistic in computer vision)  
complexity = 200 trillion, long time even on best GPU


# ANN: training

- Training set of  $m$  examples  $x^{(1)}, \dots, x^{(m)}$  (each a vector in dimension  $n_0$ ) with corresponding labels  $y^{(1)}, \dots, y^{(m)} \in \mathbb{R}$
- Learn parameters:  $b^{[1]} \in \mathbb{R}^{n_1}, b^{[2]} \in \mathbb{R}^{n_2=1}$  and  $W^{[1]} \in \mathbb{R}^{n_1 \times n_0}, W^{[2]} \in \mathbb{R}^{(n_2=1) \times n_1}$  that is  $D = n_1 + n_2 + n_1 n_0 + n_2 n_1$  parameters
- By minimizing the cost function  $C(b^{[1]}, b^{[2]}, W^{[1]}, W^{[2]}) = \sum_{i=1}^m c_i(b^{[1]}, b^{[2]}, W^{[1]}, W^{[2]})$  with loss  $c_i(b^{[1]}, b^{[2]}, W^{[1]}, W^{[2]}) = \left( a_{=\hat{y}^{(i)}}^{[2](i)} - y^{(i)} \right)^2$
- No more analytic solution ( $C$  is not even convex in the  $D$  parameter space...)
- We must apply a numerical minimization algorithm
- Many iterative algorithms exist in literature and standard libraries
  - Gradient descent and variants
  - Pseudo-Newton methods based on an estimation of the hessian (second order derivatives) matrix
  - Levenberg and Marquard's algorithm
  - And many others, see **Numerical Recipes**
- Most algorithms (and all efficient ones) require differentials of cost function to all the parameters
- On each iteration, we must compute not only the cost function but also its  $D$  derivatives where we recall that  $D \sim O(n_1 n_0)$

# Universal representation theorem

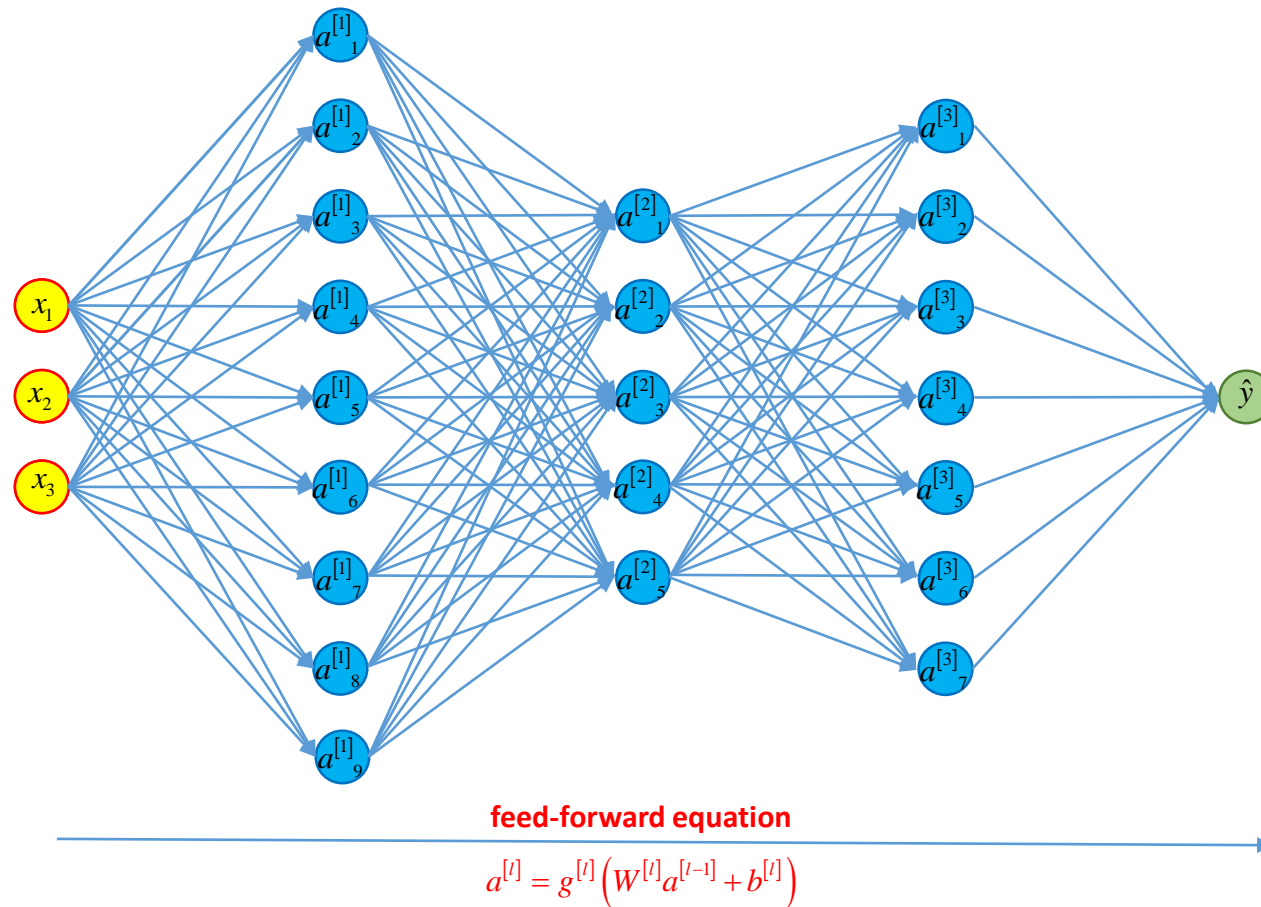
- Polynomial regression offers mathematical guarantees
  - Combinations of monomials can approximate all smooth functions to arbitrary accuracy (with max degree  $p$  high enough)
  - Hence, polynomial regression can capture (with enough monomials) can capture any (smooth) relation in the data
  - (At the cost of a high number of basis functions and the requirement of ten times that much data)
- Similarly, feed-forward ANNs with “perceptron” basis functions offer mathematical guarantees
  - Known as Universal Representation Theorem
  - States that a feed-forward perceptron may approximate any smooth function to arbitrary accuracy (with enough hidden units)
  - (Provided technical assumptions on activation functions, which sigmoidal and variants all pass)
  - See sketch of proof (with ReLu activation in dimension 1):  
<https://www.quora.com/In-Machine-Learning-ReLU-activation-function-is-said-to-be-a-universal-function-approximation-can-you-show-just-how-it-is-possible-to-approximate-any-continuous-function-by-purely-combining-many-instances-of-ReLU/answer/Antoine-Savine-1>
- In conclusion
  - ANNs are superior to basis function regression in what they learn basis functions from the data
  - MLPs offer the same guarantees as basis function regression, also computationally friendly (matrix operations fit for parallel proc)
  - But they lose analytic solution for training and require iterative cost minimization
  - Iterative minimization procedures require gradients (all differentials) of the cost function on every iteration
  - Therefore a quick computation of derivatives is key to training in reasonable time

# More on regression, basis functions and ANNs





# Deep learning: multiple hidden layers





# Deep learning: composing basis functions

- Shallow (one hidden layer) ANN:

- Prediction is a linear combination of hidden layer activations:
- Hidden layer activations are basis functions of the input layer:
- So we linearly regress on (learnt) basis functions of inputs

$$\hat{y} = W^{[2]}a^{[1]} + b^{[2]}$$
$$a^{[1]} = g^{[1]}(W^{[1]}x + b^{[1]})$$

- ANN with two hidden layers:

- Prediction is a linear combination of the **second** hidden layer:
- Second layer activations are basis functions of the first hidden layer:
- First hidden layer activations are basis functions of the input layer:
- So we regress on **basis functions of basis functions**

$$\hat{y} = W^{[3]}a^{[2]} + b^{[3]}$$
$$a^{[2]} = g^{[2]}(W^{[2]}a^{[1]} + b^{[2]})$$
$$a^{[1]} = g^{[1]}(W^{[1]}x + b^{[1]})$$

- Deep ANN with L layers (L-1 hidden layers)



- Prediction is a linear combination of the **last** hidden layer:
- Last hidden layer contains basis functions of basis functions of basis functions ... of basis functions of the inputs from the feed-forward equation:
- So we regress on many **layers of composition** of basis functions

$$\hat{y} = W^{[L]}a^{[L-1]} + b^{[L]}$$
$$a^{[l]} = g^{[l]}(W^{[l]}a^{[l-1]} + b^{[l]})$$
$$a^{[0]} = x$$

# Why deep learning?

- We know from the Universal Representation Theorem that:
  - An ANN with a single hidden layer can approximate any (smooth) function given enough units
  - So what is the point of multiple layers?
- Composition of basis function offers (exponentially) better function representation abilities
  - With a limited number of units, a single layer ANN only correctly approximates a limited sub-space of functions
  - But if we allow basis functions to compose (and learn composition from data) we can represent an exponentially wider variety of functions
  - This result should be intuitive, it was demonstrated in particular cases and widely illustrated and empirically validated in a vast number of contexts
  - In short, 10 layers of 10 units each approximate a (much) wider sup-space of functions than 1 layer with 100 units
- This is the basis for deep learning
- Deep ANNs are otherwise a trivial generalization of shallow ANNs

# Deep feed-forward networks



feed-forward equation

$$a^{[l]} = g^{[l]}(W^{[l]}a^{[l-1]} + b^{[l]})$$

- Prediction: feed-forward equations

$$a^{[0]} = x, a^{[l]} = g^{[l]}(W^{[l]}a^{[l-1]} + b^{[l]}), y = a^{[L]}$$

- Parameters: for  $1 \leq l \leq L$  :  $W^{[l]}(n_l \times n_{l-1}), b^{[l]}(n_l)$

- Number of parameters:  $D = \sum_{l=1}^L n_l(1 + n_{l-1})$

- Complexity  $\sim D = \sum_{l=1}^L n_l n_{l-1}$

- Training: find all  $W^{[l]}$  and  $b^{[l]}$  to minimise cost

$$C = \sum_{i=1}^m c_i, c_i = (\hat{y}^{(i)} - y^{(i)})^2$$

- Use iterative algorithms  
compute all  $D$  differentials on each iteration

- Key: quickly compute  
large number  $D$  of differentials

# Summary and notations: deep ANN

- Prediction

- Start with layer 0 = inputs  $\begin{bmatrix} a^{[0]}_1 \\ \dots \\ a^{[0]}_{n_0} \end{bmatrix} = \begin{bmatrix} x_1 \\ \dots \\ x_{n_0} \end{bmatrix}$
- Feed-forward from layer l-1 to layer l  $\rightarrow \begin{bmatrix} a^{[l]}_1 \\ \dots \\ a^{[l]}_{n_l} \end{bmatrix} = \phi^{[l]} \left( \begin{bmatrix} a^{[l-1]}_1 \\ \dots \\ a^{[l-1]}_{n_{l-1}} \end{bmatrix}; \begin{bmatrix} g^{[l]}_1 \\ \dots \\ g^{[l]}_{D_l} \end{bmatrix} \right)$
- Prediction = layer L  $\hat{y} = a^{[L]}_1$

- Parameters:  $\begin{bmatrix} g^{[1]}_1 \\ \dots \\ g^{[1]}_{D_1} \end{bmatrix}, \begin{bmatrix} g^{[2]}_1 \\ \dots \\ g^{[2]}_{D_2} \end{bmatrix}, \dots, \begin{bmatrix} g^{[L-1]}_1 \\ \dots \\ g^{[L-1]}_{D_{L-1}} \end{bmatrix}, \begin{bmatrix} g^{[L]}_1 \\ \dots \\ g^{[L]}_{D_L} \end{bmatrix}$  so, ultimately:  $\hat{y} = f \left( \begin{bmatrix} x_1 \\ \dots \\ x_{n_0} \end{bmatrix}; \begin{bmatrix} g^{[1]}_1 \\ \dots \\ g^{[1]}_{D_1} \end{bmatrix}, \begin{bmatrix} g^{[2]}_1 \\ \dots \\ g^{[2]}_{D_2} \end{bmatrix}, \dots, \begin{bmatrix} g^{[L-1]}_1 \\ \dots \\ g^{[L-1]}_{D_{L-1}} \end{bmatrix}, \begin{bmatrix} g^{[L]}_1 \\ \dots \\ g^{[L]}_{D_L} \end{bmatrix} \right)$

- Training: find all the parameters theta

- To minimize sum of errors on training set with known labels  $\min_{\begin{bmatrix} g^{[1]}_1 \\ \dots \\ g^{[1]}_{D_1} \end{bmatrix}, \begin{bmatrix} g^{[2]}_1 \\ \dots \\ g^{[2]}_{D_2} \end{bmatrix}, \dots, \begin{bmatrix} g^{[L-1]}_1 \\ \dots \\ g^{[L-1]}_{D_{L-1}} \end{bmatrix}, \begin{bmatrix} g^{[L]}_1 \\ \dots \\ g^{[L]}_{D_L} \end{bmatrix}} C = \sum_{i=1}^m c_i = (\hat{y}_i - y_i)^2$
- Takes iterative algorithm where all the differentials must be computed on every iteration:  $\frac{\partial c_i}{\partial g^{[l]}_j}$  for all i, l and j

# Summary and notations: MLP

- Special case of ANN with mathematical and computational benefits – by far the most common form of ANN

- Start with layer 0 = inputs 
$$\begin{bmatrix} a_1^{[0]} \\ \vdots \\ a_{n_0}^{[0]} \end{bmatrix} = \begin{bmatrix} x_1 \\ \vdots \\ x_{n_0} \end{bmatrix}$$

- Feed-forward from layer l-1 to layer l**  $\rightarrow \begin{bmatrix} a_1^{[l]} \\ \vdots \\ a_{n_l}^{[l]} \end{bmatrix} = g^{[l]} \left( \begin{bmatrix} W_{1,1}^{[l]} & \dots & W_{1,n_{l-1}}^{[l]} \\ \vdots & \ddots & \vdots \\ W_{n_l,1}^{[l]} & \dots & W_{n_l,n_{l-1}}^{[l]} \end{bmatrix} \begin{bmatrix} a_1^{[l-1]} \\ \vdots \\ a_{n_{l-1}}^{[l-1]} \end{bmatrix} + \begin{bmatrix} b_1^{[l]} \\ \vdots \\ b_{n_l}^{[l]} \end{bmatrix} \right)$  where g is scalar, applied element-wise

- Prediction = layer L  $\hat{y} = a_1^{[L]}$

- Parameters: all the  $\begin{bmatrix} W_{1,1}^{[l]} & \dots & W_{1,n_{l-1}}^{[l]} \\ \vdots & \ddots & \vdots \\ W_{n_l,1}^{[l]} & \dots & W_{n_l,n_{l-1}}^{[l]} \end{bmatrix}$  and all the  $\begin{bmatrix} b_1^{[l]} \\ \vdots \\ b_{n_l}^{[l]} \end{bmatrix}$  for all layers l

- Training: find all the parameters W and b

- Minimize sum of errors on training set with known labels 
$$\min_{\begin{bmatrix} W_{1,1}^{[l]} & \dots & W_{1,n_{l-1}}^{[l]} \\ \vdots & \ddots & \vdots \\ W_{n_l,1}^{[l]} & \dots & W_{n_l,n_{l-1}}^{[l]} \end{bmatrix}, \begin{bmatrix} b_1^{[l]} \\ \vdots \\ b_{n_l}^{[l]} \end{bmatrix}, 1 \leq l \leq L} C = \sum_{i=1}^m c_i = (\hat{y}_i - y_i)^2$$

- Takes iterative algorithm where all the differentials must be computed on every iteration:  $\frac{\partial c_i}{\partial W_{j,k}^{[l]}}, \frac{\partial c_i}{\partial b_j^{[l]}}$  for all i, l, j and k

# Back-Propagation in deep neural nets

# Differentials of the cost function

- Mutli-layer perceptrons (MLPs) predict with the feed-forward equations:
  - Recall:  $a_0 = x$  ,  $a^{[l]} = g^{[l]}(W^{[l]}a^{[l-1]} + b^{[l]})$  ,  $\hat{y} = a^{[L]}$
  - Parameters: for  $1 \leq l \leq L$  :  $W^{[l]}(n_l \times n_{l-1})$ ,  $b^{[l]}(n_l)$  ,  $D = \sum_{l=1}^L n_l(1 + n_{l-1}) \sim Ln^2$  parameters
  - Cost:  $C = \sum_{i=1}^m c_i$  ,  $c_i = (\hat{y}^{(i)} - y^{(i)})^2$
  - Must compute  $mD$  differentials  $\frac{\partial c_i}{\partial W_{j,k}^{[l]}}$ ,  $\frac{\partial c_i}{\partial b_j^{[l]}}$  on each iteration
- Feed-forward complexity:
  - For a “reasonably large” network with 10 layers of 100 units,  $D=100,000$
  - To make one prediction: complexity  $\sim D \sim 100,000$
  - We need at least  $m=10D$  training examples to avoid overfitting, so to compute one cost, complexity  $> 10D^2 \sim 10^{11}$
  - With linear differentials algorithms, we compute the cost  $D$  times to calculate  $D$  differentials, complexity  $> 10D^3 \sim 10^{16}$
  - Not viable, even in parallel on most powerful GPUs

# Differentials by finite differences

- Finite difference (FD) algorithm:
  - First compute the cost  $C_0$
  - To compute the sensitivity of the cost to each (scalar) parameter  $\theta_d, 1 \leq d \leq D$  :
    - Bump the parameter by a small amount  $\epsilon$
    - Re-compute the cost  $C_1$  , repeating feed-forward equations for all training examples
    - $\partial C / \partial \theta_d \approx (C_1 - C_0) / \epsilon$
- FD is (obviously) linear in  $D$ , therefore not viable
- But it has very desirable properties



# FD automation

- Implementation is straightforward with little scope for error
- Importantly, FD differentiation is **automatic**
  - Developers only write prediction (feed-forward) code
  - FD computes differentials automatically by calling the feed-forward code repeatedly
  - Developers don't need to write any differentiation code
- More importantly, FD automatically **synchronises** with modifications to feed-forward code
  - Tricks of the trade (see e.g. Coursera's Deep Learning Spec) to train deep nets faster and better:
    - Regularization: add a regularization term to cost function
    - Dropout: randomly drop connections between units
    - Batch Norm: normalize the mean and variance on all layers
    - And many more
  - All these modify feed-forward equations and code, and affect differentials
  - With manual differentiation, developers must ensure consistency of evaluation and differentiation code
  - This is painful and prone to error
  - With automatic differentiation, differentiation always remains consistent with evaluation
- This being said, FD and other linear algorithms are not viable for training deep nets

# Computing cost differentials: general ANN

- First notice:  $C = \sum_{i=1}^m c_i$  so  $\frac{\partial C}{\partial g^{[l]}} = \sum_{i=1}^m \frac{\partial c_i}{\partial g^{[l]}}$
- Then:  $c = (\hat{y} - y)^2 = (a^{[L]} - y)^2$  (dropping the example index  $i$  to simplify notations) so  $\frac{\partial c}{\partial \hat{y}} = \frac{\partial c}{\partial a^{[L]}} = 2(a^{[L]} - y) = \text{"2 prediction errors" (eq1)}$
- From the feed-forward equation:  $\begin{bmatrix} a^{[l]}_1 \\ \dots \\ a^{[l]}_{n_l} \end{bmatrix} = \varphi^{[l]} \left( \begin{bmatrix} a^{[l-1]}_1 \\ \dots \\ a^{[l-1]}_{n_{l-1}} \end{bmatrix}; \begin{bmatrix} g^{[l]}_1 \\ \dots \\ g^{[l]}_{D_l} \end{bmatrix} \right)$  where the functions phi and their differentials are known
- We get the Jacobian matrix of  $a^{[l]}$  to  $a^{[l-1]}$  (eq2)  $\rightarrow \frac{\partial a^{[l]}}{\partial a^{[l-1]}} = \frac{\partial \varphi^{[l]}}{\partial a^{[l-1]}} = \begin{bmatrix} \frac{\partial a^{[l]}_1}{\partial a^{[l-1]}_1} & \dots & \frac{\partial a^{[l]}_1}{\partial a^{[l-1]}_{n_{l-1}}} \\ \dots & \dots & \dots \\ \frac{\partial a^{[l]}_{n_l}}{\partial a^{[l-1]}_1} & \dots & \frac{\partial a^{[l]}_{n_l}}{\partial a^{[l-1]}_{n_{l-1}}} \end{bmatrix}$
- And from the chain rule we have eq3:  $\frac{\partial c}{\partial g^{[l]}} = \frac{\partial c}{\partial a^{[l]}} \frac{\partial a^{[l]}}{\partial g^{[l]}} = \frac{\partial c}{\partial a^{[l]}} \frac{\partial \varphi^{[l]}}{\partial g^{[l]}}$

# Computing cost differentials: strategy

- Eq1 gives  $\frac{\partial c}{\partial a^{[L]}}$  and Eq2 gives  $\frac{\partial a^{[l]}}{\partial a^{[l-1]}}$
- So we can compute for every layer l:  $\frac{\partial c}{\partial a^{[l]}} = \frac{\partial c}{\partial a^{[L]}} \frac{\partial a^{[L]}}{\partial a^{[L-1]}} \frac{\partial a^{[L-1]}}{\partial a^{[L-2]}} \cdots \frac{\partial a^{[l+1]}}{\partial a^{[l]}}$  (**back-prop equation**)
- Then, we use Eq3 to get  $\frac{\partial c}{\partial g^{[l]}} = \frac{\partial c}{\partial a^{[l]}} \frac{\partial \phi^{[l]}}{\partial g^{[l]}}$
- Back-prop equation for layer 0 (recall the output layer is a real number):

Forward computation, in the order of the layers: 0 to L

$$\frac{\partial c}{\partial a^{[0]}} = \frac{\partial c}{\partial a^{[L]}} \frac{\partial a^{[L]}}{\partial a^{[L-1]}} \frac{\partial a^{[L-1]}}{\partial a^{[L-2]}} \cdots \frac{\partial a^{[2]}}{\partial a^{[1]}} \frac{\partial a^{[1]}}{\partial a^{[0]}} = \underbrace{2(a^{[L]} - y)}_{\frac{\partial c}{\partial a^{[L]}}: \text{real number}} \underbrace{\begin{bmatrix} \frac{\partial a^{[L]}}{\partial a^{[L-1]}_1} & \cdots & \frac{\partial a^{[L]}}{\partial a^{[L-1]}_{n_{L-1}}} \end{bmatrix}}_{\frac{\partial a^{[L]}}{\partial a^{[L-1]}}: \text{vector}} \underbrace{\begin{bmatrix} \frac{\partial a^{[L-1]}_1}{\partial a^{[L-2]}_1} & \cdots & \frac{\partial a^{[L-1]}_1}{\partial a^{[L-2]}_{n_{L-2}}} \\ \cdots & \cdots & \cdots \\ \frac{\partial a^{[L-1]}_{n_{L-1}}}{\partial a^{[L-2]}_1} & \cdots & \frac{\partial a^{[L-1]}_{n_{L-1}}}{\partial a^{[L-2]}_{n_{L-2}}} \end{bmatrix}}_{\frac{\partial a^{[L-1]}}{\partial a^{[L-2]}}: \text{matrices}} \cdots \underbrace{\begin{bmatrix} \frac{\partial a^{[2]}_1}{\partial a^{[1]}_1} & \cdots & \frac{\partial a^{[2]}_1}{\partial a^{[1]}_{n_1}} \\ \cdots & \cdots & \cdots \\ \frac{\partial a^{[2]}_{n_2}}{\partial a^{[1]}_1} & \cdots & \frac{\partial a^{[2]}_{n_2}}{\partial a^{[1]}_{n_1}} \end{bmatrix}}_{\frac{\partial a^{[2]}}{\partial a^{[1]}}: \text{matrices}} \underbrace{\begin{bmatrix} \frac{\partial a^{[1]}_1}{\partial a^{[0]}_1} & \cdots & \frac{\partial a^{[1]}_1}{\partial a^{[0]}_{n_0}} \\ \cdots & \cdots & \cdots \\ \frac{\partial a^{[1]}_{n_1}}{\partial a^{[0]}_1} & \cdots & \frac{\partial a^{[1]}_{n_1}}{\partial a^{[0]}_{n_0}} \end{bmatrix}}_{\frac{\partial a^{[1]}}{\partial a^{[0]}}: \text{matrices}}$$

Backward computation, in the reverse order: L to 0

# Backprop equation: forward evaluation

- Start with the Jacobian matrix of first hidden to input layer:  $\frac{\partial a^{[1]}}{\partial a^{[0]}} = \begin{bmatrix} \frac{\partial a^{[1]}_1}{\partial a^{[0]}_1} & \dots & \frac{\partial a^{[1]}_1}{\partial a^{[0]}_{n_0}} \\ \dots & \dots & \dots \\ \frac{\partial a^{[1]}_{n_1}}{\partial a^{[0]}_1} & \dots & \frac{\partial a^{[1]}_{n_1}}{\partial a^{[0]}_{n_0}} \end{bmatrix}$
- Compute the Jacobian of second hidden to input layer  $\rightarrow \frac{\partial a^{[2]}}{\partial a^{[0]}} = \frac{\partial a^{[2]}}{\partial a^{[1]}} \frac{\partial a^{[1]}}{\partial a^{[0]}} = \begin{bmatrix} \frac{\partial a^{[2]}_1}{\partial a^{[1]}_1} & \dots & \frac{\partial a^{[2]}_1}{\partial a^{[1]}_{n_1}} \\ \dots & \dots & \dots \\ \frac{\partial a^{[2]}_{n_2}}{\partial a^{[1]}_1} & \dots & \frac{\partial a^{[2]}_{n_2}}{\partial a^{[1]}_{n_1}} \end{bmatrix} \begin{bmatrix} \frac{\partial a^{[1]}_1}{\partial a^{[0]}_1} & \dots & \frac{\partial a^{[1]}_1}{\partial a^{[0]}_{n_0}} \\ \dots & \dots & \dots \\ \frac{\partial a^{[1]}_{n_1}}{\partial a^{[0]}_1} & \dots & \frac{\partial a^{[1]}_{n_1}}{\partial a^{[0]}_{n_0}} \end{bmatrix} = \begin{bmatrix} \frac{\partial a^{[2]}_1}{\partial a^{[0]}_1} & \dots & \frac{\partial a^{[2]}_1}{\partial a^{[0]}_{n_0}} \\ \dots & \dots & \dots \\ \frac{\partial a^{[2]}_{n_2}}{\partial a^{[0]}_1} & \dots & \frac{\partial a^{[2]}_{n_2}}{\partial a^{[0]}_{n_0}} \end{bmatrix}$
- Carry on multiplying matrices up to layer L-1 to sequentially obtain the Jacobians matrices of hidden layer 3, 4, ..., L-1 to input layer:
 
$$\frac{\partial a^{[L-1]}}{\partial a^{[0]}} = \frac{\partial a^{[L-1]}}{\partial a^{[L-2]}} \frac{\partial a^{[L-2]}}{\partial a^{[0]}} = \begin{bmatrix} \frac{\partial a^{[L-1]}_1}{\partial a^{[L-2]}_1} & \dots & \frac{\partial a^{[L-1]}_1}{\partial a^{[L-2]}_{n_{L-2}}} \\ \dots & \dots & \dots \\ \frac{\partial a^{[L-1]}_{n_{L-1}}}{\partial a^{[L-2]}_1} & \dots & \frac{\partial a^{[L-1]}_{n_{L-1}}}{\partial a^{[L-2]}_{n_{L-2}}} \end{bmatrix} \begin{bmatrix} \frac{\partial a^{[L-2]}_1}{\partial a^{[0]}_1} & \dots & \frac{\partial a^{[L-2]}_{n_{L-2}}}{\partial a^{[0]}_{n_0}} \\ \dots & \dots & \dots \\ \frac{\partial a^{[L-2]}_1}{\partial a^{[0]}_1} & \dots & \frac{\partial a^{[L-2]}_{n_{L-2}}}{\partial a^{[0]}_{n_0}} \end{bmatrix} = \begin{bmatrix} \frac{\partial a^{[L-1]}_1}{\partial a^{[0]}_1} & \dots & \frac{\partial a^{[L-1]}_{n_{L-1}}}{\partial a^{[0]}_{n_0}} \\ \dots & \dots & \dots \\ \frac{\partial a^{[L-1]}_1}{\partial a^{[0]}_1} & \dots & \frac{\partial a^{[L-1]}_{n_{L-1}}}{\partial a^{[0]}_{n_0}} \end{bmatrix}$$
- Finally, compute the gradient of output to input layer:  $\frac{\partial a^{[L]}}{\partial a^{[0]}} = \frac{\partial a^{[L]}}{\partial a^{[L-1]}} \frac{\partial a^{[L-1]}}{\partial a^{[0]}} = \begin{bmatrix} \frac{\partial a^{[L]}}{\partial a^{[L-1]}_1} & \dots & \frac{\partial a^{[L]}}{\partial a^{[L-1]}_{n_{L-1}}} \\ \dots & \dots & \dots \\ \frac{\partial a^{[L]}}{\partial a^{[L-1]}_1} & \dots & \frac{\partial a^{[L]}}{\partial a^{[L-1]}_{n_{L-1}}} \end{bmatrix} \begin{bmatrix} \frac{\partial a^{[L-1]}_1}{\partial a^{[0]}_1} & \dots & \frac{\partial a^{[L-1]}_{n_{L-1}}}{\partial a^{[0]}_{n_0}} \\ \dots & \dots & \dots \\ \frac{\partial a^{[L-1]}_1}{\partial a^{[0]}_1} & \dots & \frac{\partial a^{[L-1]}_{n_{L-1}}}{\partial a^{[0]}_{n_0}} \end{bmatrix} = \begin{bmatrix} \frac{\partial a^{[L]}}{\partial a^{[0]}_1} & \dots & \frac{\partial a^{[L]}}{\partial a^{[0]}_{n_0}} \end{bmatrix}$
- And the gradient of the loss:  $\frac{\partial c}{\partial a^{[0]}} = \frac{\partial c}{\partial a^{[L]}} \frac{\partial a^{[L]}}{\partial a^{[0]}} = 2(a^{[L]} - y) \begin{bmatrix} \frac{\partial a^{[L]}}{\partial a^{[0]}_1} & \dots & \frac{\partial a^{[L]}}{\partial a^{[0]}_{n_0}} \end{bmatrix}$

# Forward evaluation: comments

- We recursively propagate **Jacobian matrices** to inputs  $\frac{\partial a^{[L]}}{\partial a^{[0]}}$  from  $\frac{\partial a^{[1]}}{\partial a^{[0]}}$  to  $\frac{\partial a^{[L]}}{\partial a^{[0]}}$
- For each layer, we multiply matrices to obtain matrices, which is a **cubic** operation
- After we have  $\frac{\partial a^{[L]}}{\partial a^{[0]}}$  and  $\frac{\partial c}{\partial a^{[0]}}$ , we must repeat the process to find  $\frac{\partial c}{\partial a^{[1]}}, \frac{\partial c}{\partial a^{[2]}}, \dots, \frac{\partial c}{\partial a^{[L-1]}}$
- The last comment is particular to ANNs (where we need gradients of loss to all layers)  
But the other comments are general (they apply whenever we need gradients of a final result to inputs)

# Backprop equation: backward evaluation

- Start with the gradient vector of loss (real number) to last hidden layer L-1 (vector): 
$$\frac{\partial c}{\partial a^{[L-1]}} = \underbrace{2(a^{[L]} - y)}_{\substack{\frac{\partial c}{\partial a^{[L]}}; \text{ real number}}} \underbrace{\begin{bmatrix} \frac{\partial a^{[L]}}{\partial a^{[L-1]}_1} & \cdots & \frac{\partial a^{[L]}}{\partial a^{[L-1]}_{n_{L-1}}} \end{bmatrix}}_{\substack{\frac{\partial a^{[L]}}{\partial a^{[L-1]}}; \text{ vector}}}$$
- Compute the gradient vector of loss to hidden layer L-2: 
$$\frac{\partial c}{\partial a^{[L-2]}} = \frac{\partial c}{\partial a^{[L-1]}} \frac{\partial a^{[L-1]}}{\partial a^{[L-2]}} = \begin{bmatrix} \frac{\partial c}{\partial a^{[L-1]}_1} & \cdots & \frac{\partial c}{\partial a^{[L-1]}_{n_{L-1}}} \end{bmatrix} \begin{bmatrix} \frac{\partial a^{[L-1]}_1}{\partial a^{[L-2]}_1} & \cdots & \frac{\partial a^{[L-1]}_1}{\partial a^{[L-2]}_{n_{L-2}}} \\ \cdots & \cdots & \cdots \\ \frac{\partial a^{[L-1]}_{n_{L-1}}}{\partial a^{[L-2]}_1} & \cdots & \frac{\partial a^{[L-1]}_{n_{L-1}}}{\partial a^{[L-2]}_{n_{L-2}}} \end{bmatrix} = \begin{bmatrix} \frac{\partial c}{\partial a^{[L-2]}_1} & \cdots & \frac{\partial c}{\partial a^{[L-2]}_{n_{L-2}}} \end{bmatrix}$$
- Carry on multiplying vectors by matrices down to layer 0 to recursively obtain the gradient vectors of loss to layer L-3, L-4, ..., 1, 0:

$$\frac{\partial c}{\partial a^{[0]}} = \frac{\partial c}{\partial a^{[1]}} \frac{\partial a^{[1]}}{\partial a^{[0]}} = \begin{bmatrix} \frac{\partial c}{\partial a^{[1]}_1} & \cdots & \frac{\partial c}{\partial a^{[1]}_{n_1}} \end{bmatrix} \begin{bmatrix} \frac{\partial a^{[1]}_1}{\partial a^{[0]}_1} & \cdots & \frac{\partial a^{[1]}_1}{\partial a^{[0]}_{n_0}} \\ \cdots & \cdots & \cdots \\ \frac{\partial a^{[1]}_{n_1}}{\partial a^{[0]}_1} & \cdots & \frac{\partial a^{[1]}_{n_1}}{\partial a^{[0]}_{n_0}} \end{bmatrix} = \begin{bmatrix} \frac{\partial c}{\partial a^{[0]}_1} & \cdots & \frac{\partial c}{\partial a^{[0]}_{n_0}} \end{bmatrix}$$

# Backward evaluation: comments

- We recursively propagate **adjoint vectors**  $\frac{\partial c}{\partial a^{[l]}}$  from  $\frac{\partial c}{\partial a^{[L-1]}}$  to  $\frac{\partial c}{\partial a^{[0]}}$
- Definition: given a result  $c$ , the **adjoint** of some number/vector/matrix  $x$  is  $\frac{\partial c}{\partial x}$
- Note: the adjoint of a scalar result  $c$  is of the dimension of  $x$ , hence  $\frac{\partial c}{\partial a^{[l]}}$  is a vector
- For each layer, we multiply vectors by matrices to obtain vectors, which is a **quadratic** operation
- For ANNs, we repeat the process only once to find all the  $\frac{\partial c}{\partial a^{[L-1]}}, \dots, \frac{\partial c}{\partial a^{[2]}}, \frac{\partial c}{\partial a^{[1]}}, \frac{\partial c}{\partial a^{[0]}}$
- It follows that backward evaluation is one order of magnitude faster than forward evaluation
  - Because the output layer is scalar, its adjoints are vectors
  - Whereas the input layer is a vector, hence, its Jacobians are matrices
  - To propagate adjoints is (one order of magnitude) more efficient!

# Computing cost differentials: MLPs

- Differentiation equations

- Eq1 is unchanged:  $c = (\hat{y} - y)^2 = (a^{[L]} - y)^2$  so  $\frac{\partial c}{\partial \hat{y}} = \frac{\partial c}{\partial a^{[L]}} = 2(a^{[L]} - y) = \text{"2 prediction errors"}$

- Recall MLPs are defined by the feed-forward equation:  $a^{[l]} = g^{[l]} \left( \underbrace{W^{[l]} a^{[l-1]} + b^{[l]}}_{z^{[l]}} \right)$  where  $g$  is applied element-wise to the vector  $z$

- By differentiation, we have the Jacobian matrix of  $a^{[l]}$  to  $a^{[l-1]}$  (eq2):  $\frac{\partial a^{[l]}}{\partial a^{[l-1]}} = \frac{\partial a^{[l]}}{\partial z^{[l]}} \frac{\partial z^{[l]}}{\partial a^{[l-1]}} = \underbrace{\text{diag} \left[ \underbrace{g^{[l]'}(z^{[l]})}_{n_l \times 1} \right]}_{n_l \times n_{l-1}} W^{[l]}$

- And derivatives to parameters (eq3):  $\frac{\partial c}{\partial W^{[l]}_{ij}} = \frac{\partial c}{\partial a^{[l]}_i} \frac{\partial a^{[l]}_i}{\partial W^{[l]}_{ij}} = \frac{\partial c}{\partial a^{[l]}_i} g^{[l]'}(z^{[l]}_i) a^{[l-1]}_j$  and  $\frac{\partial c}{\partial b^{[l]}_i} = \frac{\partial c}{\partial a^{[l]}_i} \frac{\partial a^{[l]}_i}{\partial b^{[l]}_i} = \frac{\partial c}{\partial a^{[l]}_i} g^{[l]'}(z^{[l]}_i)$

- Differentiation strategy

- Apply eq1 and eq2 to compute the adjoints  $\frac{\partial c}{\partial a^{[l]}}$  for all layers  $l$
- Then, apply eq3 to compute the desired sensitivities to parameters  $\frac{\partial c}{\partial W^{[l]}_{ij}}$  and  $\frac{\partial c}{\partial b^{[l]}_i}$



# Jacobian propagation

- Computation of  $\frac{\partial c}{\partial a^{[l]}}$  by Jacobian propagation in increasing order of layers:

- Start with the Jacobian matrix  $\frac{\partial a^{[l+1]}}{\partial a^{[l]}}$ , known from eq2

- Multiply by the Jacobian matrix  $\frac{\partial a^{[l+2]}}{\partial a^{[l+1]}}$  to produce the Jacobian matrix  $\frac{\partial a^{[l+2]}}{\partial a^{[l]}} = \frac{\partial a^{[l+2]}}{\partial a^{[l+1]}} \frac{\partial a^{[l+1]}}{\partial a^{[l]}}$

- Perform  $L-l$  matrix by matrix products (cubic complexity) to obtain Jacobian matrices  $\frac{\partial a^{[L]}}{\partial a^{[l]}} = \frac{\partial a^{[L]}}{\partial a^{[L-1]}} \frac{\partial a^{[L-1]}}{\partial a^{[L-2]}} \cdots \frac{\partial a^{[l+1]}}{\partial a^{[l]}}$

- Finally,  $\frac{\partial c}{\partial a^{[l]}} = \frac{\partial c}{\partial a^{[L]}} \frac{\partial a^{[L]}}{\partial a^{[l]}}$  where  $\frac{\partial c}{\partial a^{[L]}}$  is known from eq1

- Repeat for every layer  $l$

$$\frac{\partial a^{[l+1]}}{\partial a^{[l]}} \rightarrow \frac{\partial a^{[l+2]}}{\partial a^{[l]}} \rightarrow \frac{\partial a^{[l+3]}}{\partial a^{[l]}} \rightarrow \frac{\partial a^{[l+k]}}{\partial a^{[l]}} \rightarrow \frac{\partial a^{[L]}}{\partial a^{[l]}} \rightarrow \frac{\partial c}{\partial a^{[l]}}$$

- Obviously inefficient:

- Must repeat for every level  $l$
- Sequence of expensive (cubic) matrix by matrix products

# Adjoint propagation

- Computation of all the  $\frac{\partial c}{\partial a^{[l]}}$  by adjoint propagation in the reverse order of layers:
    - Start with the **real number**  $\frac{\partial c}{\partial a^{[L]}}$  known from eq1
    - Multiply by the **gradient vector**  $\frac{\partial a^{[L]}}{\partial a^{[L-1]}}$  known from eq2 to produce the **adjoint vector**  $\frac{\partial c}{\partial a^{[L-1]}} = \frac{\partial c}{\partial a^{[L]}} \frac{\partial a^{[L]}}{\partial a^{[L-1]}}$
    - Perform  $L$  **matrix by vector** products  $\frac{\partial c}{\partial a^{[l-1]}} = \frac{\partial c}{\partial a^{[l]}} \frac{\partial a^{[l]}}{\partial a^{[l-1]}}$  (quadratic complexity) to obtain all the adjoint vectors  $\frac{\partial c}{\partial a^{[l]}}$
- $$\frac{\partial c}{\partial a^{[1]}} \leftarrow \frac{\partial c}{\partial a^{[2]}} \leftarrow \frac{\partial c}{\partial a^{[3]}} \leftarrow \frac{\partial c}{\partial a^{[4]}} \leftarrow \frac{\partial c}{\partial a^{[5]}} \leftarrow \frac{\partial c}{\partial a^{[6]}}$$
- Complexity: all  $\frac{\partial c}{\partial a^{[l]}}$  and hence all  $\frac{\partial c}{\partial W^{[l]}_{ij}}$  and  $\frac{\partial c}{\partial b^{[l]}_i}$  are computed with  $L$  matrix by vector products
    - Same complexity as for making a prediction / computing a loss
    - Constant time differentiation:** all differentials of  $c$  are produced with same (leading) complexity as one evaluation of  $c$

# Back-Propagation

- Reverse adjoint propagation
  - Also called back-propagation or back-prop in ML lingo, or **adjoint differentiation** (AD) in general
  - Computes all differentials of the cost to the many parameters in a deep net in a time similar to one evaluation of the cost
  - Achieves this remarkable result by reversing the order of the calculations in the differentiation
  - Relies on the cost being scalar, hence its adjoints are vectors
  - Is the only viable means to train deep nets in reasonable time
  - Is the reason why deep learning achieved such spectacular success and why your phone quickly learns to recognise you
- However manual AD code is
  - Complicated and prone to error
  - Painful to maintain and synchronize as we modify feed-forward code
- AAD: **automatic** adjoint differentiation
  - Next, we learn to **automate** AD
  - Since AD propagate adjoints in the reverse order through evaluation graphs
  - We will learn to automatically generate graphs
  - Not only for deep nets, but for any calculation
  - Including a large Monte-Carlo simulation

AD through evaluation graphs

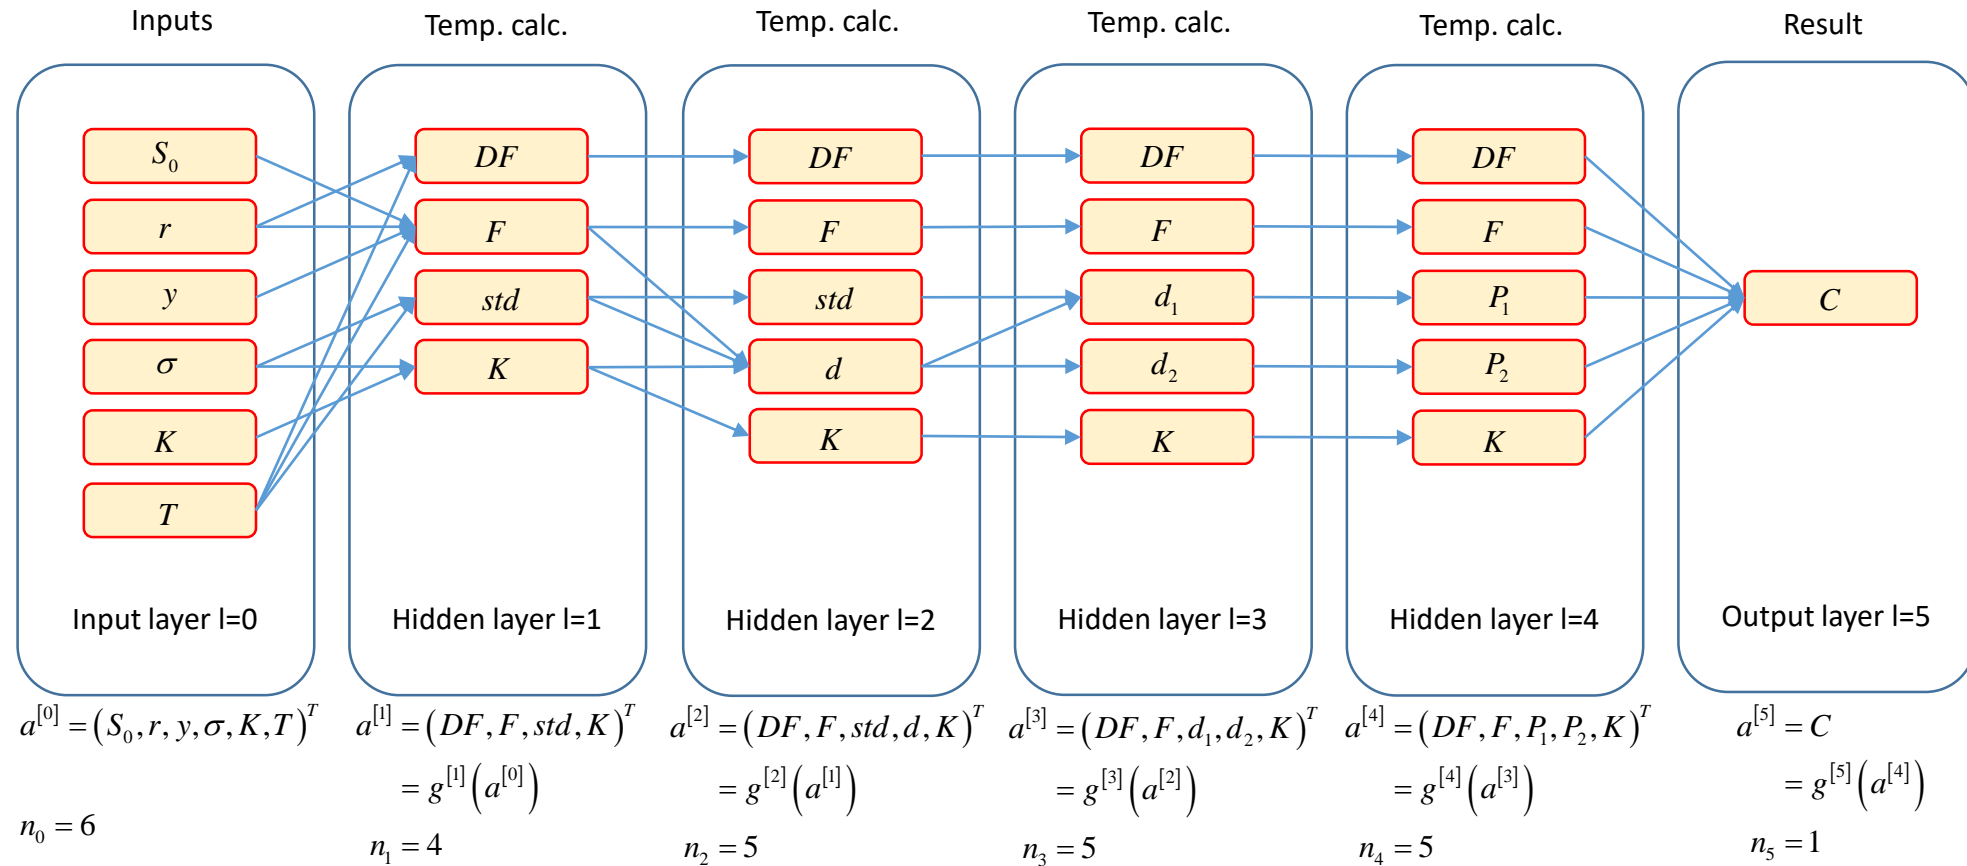
# Evaluation graphs and adjoint propagation

- Evaluation graph
  - We have seen evaluation graphs for neural nets
  - Evaluation graphs express the order and dependencies of “atomic” computations involved in a calculation
  - Any calculation defines an evaluation graph
  - To some extent, every calculation defines a sequence of feed-forward operations, similarly to a neural net
  - Chapter 9 of the AAD book teaches to build evaluation graphs in memory with operator overloading in C++
- Adjoint Differentiation
  - Is not limited to neural nets
  - Is applicable to any calculation
  - Propagates adjoints through the evaluation graph
  - Offers constant time efficiency (all differentials for an expense similar to one evaluation) for scalar calculations
  - (A scalar calculation is one where the final result –or output layer, is scalar)
- Once we have an evaluation graph, AD may be automated

# Example: Black & Scholes

- Black & Scholes' formula:  $C(S_0, r, y, \sigma, K, T) = DF [FN(d_1) - KN(d_2)]$  with
  - Discount factor to maturity:  $DF = \exp(-rT)$
  - Forward:  $F = S_0 \exp[(r - y)T]$
  - Standard deviation:  $std = \sigma\sqrt{T}$
  - Log-moneyness:  $d = \frac{\log\left(\frac{F}{K}\right)}{std}$
  - D's:  $d_1 = d + \frac{std}{2}$ ,  $d_2 = d - \frac{std}{2}$
  - Probabilities to end in the money, resp. under spot and risk-neutral measures:  $P_1 = N(d_1)$ ,  $P_2 = N(d_2)$
  - Call price:  $C = DF [FP_1 - KP_2]$

# Black & Scholes: evaluation graph



# Feed-forward equations

- As in deep nets, for **any** calculation, we have the feed-forward equation defining the evaluation graph:

inputs:  $a^{[0]} = x$

feed-forward:  $a^{[l]} = g^{[l]}(a^{[l-1]})$

result:  $y = a^{[L]}$

- In Black & Scholes, we have

$$g^{[1]} \begin{pmatrix} x_1 \\ x_2 \\ x_3 \\ x_4 \\ x_5 \\ x_6 \end{pmatrix} = \begin{pmatrix} \exp(-x_2 x_6) \\ x_1 \exp[(x_2 - x_3) x_6] \\ x_4 \sqrt{x_6} \\ x_5 \end{pmatrix} \quad g^{[2]} \begin{pmatrix} x_1 \\ x_2 \\ x_3 \\ x_4 \end{pmatrix} = \begin{pmatrix} x_1 \\ x_2 \\ x_3 \\ \frac{\log(x_2/x_4)}{x_3} \\ x_4 \end{pmatrix} \quad g^{[3]} \begin{pmatrix} x_1 \\ x_2 \\ x_3 \\ x_4 \\ x_5 \end{pmatrix} = \begin{pmatrix} x_1 \\ x_2 \\ x_4 + \frac{x_3}{2} \\ x_4 - \frac{x_3}{2} \\ x_5 \end{pmatrix} \quad g^{[4]} \begin{pmatrix} x_1 \\ x_2 \\ x_3 \\ x_4 \\ x_5 \end{pmatrix} = \begin{pmatrix} x_1 \\ x_2 \\ N(x_3) \\ N(x_4) \\ x_5 \end{pmatrix} \quad g^{[5]} \begin{pmatrix} x_1 \\ x_2 \\ x_3 \\ x_4 \\ x_5 \end{pmatrix} = x_1 (x_2 x_3 - x_5 x_4)$$



# Differentiation equations

- For **any** calculation, where we compute the differentials of the result to the inputs:  $\frac{\partial y}{\partial x} = \frac{\partial a^{[L]}}{\partial a^{[0]}}$
- If we split the calculation into small enough pieces, we can easily compute the Jacobians  $g^{[l]'} = \frac{\partial a^{[l]}}{\partial a^{[l-1]}}$
- In Black & Scholes

$$g^{[1]'} = \begin{pmatrix} 0 & -x_6 \exp(-x_2 x_6) & 0 & 0 & 0 & -x_2 \exp(-x_2 x_6) \\ \exp[(x_2 - x_3) x_6] & x_1 x_6 \exp[(x_2 - x_3) x_6] & -x_1 x_6 \exp[(x_2 - x_3) x_6] & 0 & 0 & x_1 (x_2 - x_3) \exp[(x_2 - x_3) x_6] \\ 0 & 0 & 0 & \sqrt{x_6} & 0 & \frac{x_4}{2\sqrt{x_6}} \\ 0 & 0 & 0 & 0 & 1 & 0 \end{pmatrix}$$

$$g^{[2]'} = \begin{pmatrix} 1 & 0 & 0 & 0 \\ 0 & 1 & 0 & 0 \\ 0 & 0 & 1 & 0 \\ 0 & \frac{1}{x_2 x_3} & -\frac{\log(x_2/x_4)}{x_3^2} & -\frac{1}{x_3 x_4} \\ 0 & 0 & 0 & 1 \end{pmatrix}$$

$$g^{[3]'} = \begin{pmatrix} 1 & 0 & 0 & 0 & 0 \\ 0 & 1 & 0 & 0 & 0 \\ 0 & 0 & \frac{1}{2} & 1 & 0 \\ 0 & 0 & -\frac{1}{2} & 1 & 0 \\ 0 & 0 & 0 & 0 & 1 \end{pmatrix}$$

$$g^{[4]'} = \begin{pmatrix} 1 & 0 & 0 & 0 & 0 \\ 0 & 1 & 0 & 0 & 0 \\ 0 & 0 & n(x_3) & 0 & 0 \\ 0 & 0 & 0 & n(x_4) & 0 \\ 0 & 0 & 0 & 0 & 1 \end{pmatrix}$$

$$g^{[5]'} = ((x_2 x_3 - x_5 x_4), x_1 x_3, x_1 x_2, -x_1 x_5, -x_1 x_4)$$

- Note that the Jacobian of the output layer is a vector, as with any scalar calculation

# Forward and backward differentiation

- For **any** calculation, we have:  $\frac{\partial y}{\partial x} = \frac{\partial a^{[L]}}{\partial a^{[0]}} = \frac{\partial a^{[L]}}{\partial a^{[L-1]}} \frac{\partial a^{[L-1]}}{\partial a^{[L-2]}} \cdots \frac{\partial a^{[l]}}{\partial a^{[l-1]}} \cdots \frac{\partial a^{[2]}}{\partial a^{[1]}} \frac{\partial a^{[1]}}{\partial a^{[0]}}$
- If we compute this product right to left:
  - We follow the order of the calculation
  - We propagate Jacobians to the inputs up from layer 1 to the output layer
  - We repeatedly multiply matrices with matrices, with cubic complexity
- If we compute this product left to right:
  - We follow the reverse order of the calculation
  - We propagate adjoints of the result down from layer  $L-1$  to the input layer
  - Provided the final result is scalar, we repeatedly multiply matrices by vectors, with quadratic complexity
- Left as an exercise:  
show that Black & Scholes' Jacobians of the previous slide lead to the well known Black-Scholes "Greeks"

# Evaluation graphs: conclusion

- All calculations define a graph
- Like deep nets, evaluation graphs may be organized in successive layers function of one another:  $a^{[l]} = g^{[l]}(a^{[l-1]})$
- Once the evaluation graph is known down to elementary operations:
  - The Jacobians  $g^{[l]'} = \frac{\partial a^{[l]}}{\partial a^{[l-1]}}$  are found trivially
  - Adjoint differentiation may be conducted an order of magnitude faster than Jacobian propagation (provided the calculation is scalar):

$$\frac{\partial y}{\partial x} = \frac{\partial a^{[L]}}{\partial a^{[0]}} = \frac{\partial a^{[L]}}{\partial a^{[L-1]}} \frac{\partial a^{[L-1]}}{\partial a^{[L-2]}} \cdots \frac{\partial a^{[l]}}{\partial a^{[l-1]}} \cdots \frac{\partial a^{[2]}}{\partial a^{[1]}} \frac{\partial a^{[1]}}{\partial a^{[0]}}$$

↗ adjoint propagation, quadratic  
↖ Jacobian propagation, cubic

- Next, we show how evaluation graphs may be **automatically** generated

Recording calculations on tape

# Code

- From there on, we work with C++ code
- We only work with simplistic, beginner level C++
- We show that we can still implement AAD with rather spectacular results
- We refer to the AAD book for professional, generic, parallel, efficient implementation
- **All the code is freely available on GitHub, in the file toyCode.h**

<http://github.com/asavine>

# Automatic differentiation

- Differentiation by finite differences
  - Computes derivatives by running evaluation repeatedly
  - Only requires executable evaluation code
  - Linear complexity in the number of differentials
- AAD
  - Computes derivatives by running AD in reverse order over an evaluation graph
  - Requires a graph to traverse – executable code is not enough
  - Constant complexity in the number of differentials : fast differentiation of scalar code
  - To implement AAD we must extract an evaluation graph: sequence of calculations (nodes) and their dependencies (edges)

# Building evaluation graphs

- Explicit evaluation graphs
  - Solution implemented in TensorFlow
  - Lets users explicitly build graphs by calling graph creation functions exported to many languages
  - Then call TensorFlow API to evaluate or differentiate the graphs efficiently: AD, parallel CPU, GPU
  - Smart and efficient
  - But forces developers to explicitly build graphs in place of calculation code
  - (TensorFlow has a nice API that makes building graphs similar to coding calculations)
  - What we want here is take some calculation code and automatically extract its graph
- Source transformation
  - Code that reads and understands evaluation code, extracts graph and writes backward differentiation code automatically
  - Complex, specialized work similar to writing compilers
  - Variant: template meta-programming and introspection (introduced in chapter 15)

# Operator overloading

- Alternative: operator overloading (in languages that support it like C++)
- When code applies operators (like + or \*) or math functions (like log or sqrt) to real numbers (*double* type) the corresponding operations are evaluated immediately (or *eagerly*):

```
double x = 1, y = 2;    // x and y are doubles
double z = x + y;       // evaluates x + y and stores result in z
double t = log(x);       // evaluates log(x) and stores result in t
```

- When the same operators are applied to **custom types** (our own type to store real numbers) developers decide what exactly is executed:

```
class myNumberType
{
    // class definition here
};

myNumberType operator+(const myNumberType& lhs, const myNumberType& rhs)
{
    // this code is executed anytime two numbers of type myNumberType are added
}

myNumberType log(const myNumberType& arg)
{
    // this code is executed anytime log is called on a number of type myNumberType
}

myNumberType x, y;      // x and y are of type myNumberType
myNumberType z = x + y;  // executes code in operator+
myNumberType t = log(x); // executes code in log
```



# Recording operations

- We apply operator overloading to **record** all operations on tape:

```
class myNumberType
{
    myNumberType(const double x)
    {
        // constructor initializes value to x and record node
    }
};

myNumberType operator+(const myNumberType& lhs, const myNumberType& rhs)
{
    recordAddition(lhs, rhs); // records addition with dependency on lhs and rhs
}

myNumberType log(const myNumberType& arg)
{
    recordLog(arg); // records log with dependency on arg
}

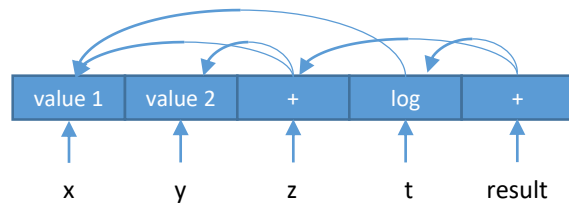
myNumberType x = 1, y = 2; // initializes x to 1 and y to 2 and records them
myNumberType z = x + y; // records addition
myNumberType t = log(x); // records log
myNumberType result = t + z; // records addition
```

# Lazy evaluation

- When this code is executed:

```
myNumberType x = 1, y = 2;    // initializes x to 1 and y to 2 and records them
myNumberType z = x + y;      // records addition
myNumberType t = log(x);     // records log
myNumberType result = t + z; // records addition
```

- Nothing is calculated, instead the following sequence is recorded in memory:



- This sequence can be evaluated later (*lazy evaluation*) or differentiated (applying AD from back to front)
- This is how we build the evaluation graph at run time, by executing calculation code with a custom type for which operators are overloaded to perform recording

# Conventional implementation

- We introduce the conventional implementation of AAD
  - Where every math operation: +, -, \*, /, pow, log, exp, sqrt, ... is recorded on a data structure called **tape**
  - (It is really an evaluation graph but we call it “tape” in AAD lingo)
- Cutting-edge implementation records whole expressions
  - With template meta-programming and expression templates
  - Resulting in 2x to 5x faster code
  - Sometimes called “tape compression”
  - Extreme instance: “tapeless AD” where we don’t record anything!
  - All implemented in professional code on GitHub and explained in detail in chapter 15
  - In this presentation, we stick with conventional implementation
- We record every mathematical operation involved in a calculation
  - Every calculation, up to complex Monte-Carlo simulations, is a sequence of +, -, \*, /, pow, log, exp, sqrt, ... !
  - We record every single one
  - **Note that all operations have either 0, 1 or 2 arguments**

# Simplistic implementation

- In this presentation, we focus on the simplicity of the code
  - We use basic C++ code and disregard efficiency, scalability and best practice
  - Our aim is to explain the key ideas with simplistic code
  - This code works, just not efficiently
- In the book, on the contrary, we build professional, scalable, efficient code in modern C++
- This is particularly important for AAD because of:
  - Recording overhead
    - Every addition, multiplication, etc. produces a record
    - Recording necessarily involves an overhead
    - An efficient implementation must minimize overhead and make recording as efficient as possible
  - Vast memory consumption
    - We store in memory all the operations involved in a large calculation, this is a very large number of records
    - Estimated RAM consumption around 5GB per second
    - Efficient memory and cache management are key to an efficient implementation
- An efficient implementation of conventional AAD is given and explained in chapter 10
  - Effective, custom memory management
  - Recording with minimum overhead, cache efficiency and so on

# Record and tape data structures

- A record
  - Stores one operation  $y = f(x)$  where  $f = +, *, -, /, \log, \sqrt{\phantom{x}}, \dots$
  - Knows the number and location of the 0, 1 or 2 arguments  $x_i$
  - Stores the 0, 1 or 2 partial derivatives to its arguments  $\partial f / \partial x_i$  so we can apply AD
- The tape stores the sequence of records

```
struct Record
{
    int    numArg;      // number of arguments: 0, 1 or 2
    int    idx1;        // index of first argument on tape
    int    idx2;        // index of second argument on tape
    double der1;        // partial derivative to first argument
    double der2;        // partial derivative to second argument
};

// The tape, declared as a global variable
vector<Record> tape;
```

# Custom real number

- Our custom number
  - Holds its value and
  - Knows the index of the corresponding operation on tape
  - May be initialized with a value to create a record (without arguments) on tape

```
struct Number
{
    double value;
    int idx;

    // default constructor does nothing
    Number() {}

    // constructs with a value and record
    Number(const double& x) : value(x)
    {
        // create a new record on tape
        tape.push_back(Record());
        Record& rec = tape.back();

        // reference record on tape
        idx = tape.size() - 1;

        // populate record on tape
        rec.numArg = 0;
    }
};
```

- We overload all mathematical operators and functions to:
  - Evaluate and store result as usual
  - Additionally, record the operation and its derivatives on tape
  - So code is evaluated and recorded at the same time

```
Number operator+(const Number& lhs, const Number& rhs)
{
    // create a new record on tape
    tape.push_back(Record());
    Record& rec = tape.back();

    // compute result
    Number result;
    result.value = lhs.value + rhs.value; // calling double overload

    // reference record on tape
    result.idx = tape.size() - 1;

    // populate record on tape
    rec.numArg = 2;
    rec.idx1 = lhs.idx;
    rec.idx2 = rhs.idx;

    // compute derivatives, both derivatives of addition are 1
    rec.der1 = 1;
    rec.der2 = 1;

    return result;
}
```

# Operator overloading

- Similarly, we overload -, \*, and /
- Same code exactly, only values and derivatives change

```
Number operator-(const Number& lhs, const Number&rhs)
{
    // ...

    // compute value
    result.value = lhs.value - rhs.value;

    // ...

    // compute derivatives
    rec.der1 = 1;
    rec.der2 = -1;

    // ...
}
```

```
Number operator*(const Number& lhs, const Number&rhs)
{
    // ...

    // compute value
    result.value = lhs.value * rhs.value;

    // ...

    // compute derivatives
    rec.der1 = rhs.value;
    rec.der2 = lhs.value;

    // ...
}
```

```
Number operator/(const Number& lhs, const Number&rhs)
{
    // ...

    // compute value
    result.value = lhs.value / rhs.value;

    // ...

    // compute derivatives
    rec.der1 = 1.0 / rhs.value;
    rec.der2 = - lhs.value / (rhs.value * rhs.value);

    // ...
}
```

# On-class operator overloading

- We must also overload +=, -=, \*= and /=, as well as unary + and -, on class

The final Number class is therefore:

```
struct Number
{
    double value;
    int idx;

    // default constructor does nothing
    Number() {}

    // constructs with a value and record
    Number(const double& x) : value(x)
    {
        // create a new record on tape
        tape.push_back(Record());
        Record& rec = tape.back();

        // reference record on tape
        idx = tape.size() - 1;

        // populate record on tape
        rec.numArg = 0;
    }

    Number operator +() const { return *this; }
    Number operator -() const { return Number(0.0) - *this; }

    Number& operator +=(const Number& rhs) { *this = *this + rhs; return *this; }
    Number& operator -=(const Number& rhs) { *this = *this - rhs; return *this; }
    Number& operator *=(const Number& rhs) { *this = *this * rhs; return *this; }
    Number& operator /=(const Number& rhs) { *this = *this / rhs; return *this; }
};
```



# Function overloading

- Similarly, we overload log, exp, sqrt
- We should really overload all standard math functions
- Code is identical for all functions, only value and derivatives change

```
Number log(const Number& arg)
{
    // create a new record on tape
    tape.push_back(Record());
    Record& rec = tape.back();

    // compute result
    Number result;
    result.value = log(arg.value);

    // reference record on tape
    result.idx = tape.size() - 1;

    // populate record on tape
    rec.numArg = 1;
    rec.idx1 = arg.idx;

    // compute derivative
    rec.der1 = 1.0 / arg.value;

    return result;
}
```

```
Number exp(const Number& arg)
{
    // create a new record on tape
    tape.push_back(Record());
    Record& rec = tape.back();

    // compute result
    Number result;
    result.value = exp(arg.value);

    // reference record on tape
    result.idx = tape.size() - 1;

    // populate record on tape
    rec.numArg = 1;
    rec.idx1 = arg.idx;

    // compute derivative
    rec.der1 = result.value;

    return result;
}
```

```
Number sqrt(const Number& arg)
{
    // create a new record on tape
    tape.push_back(Record());
    Record& rec = tape.back();

    // compute result
    Number result;
    result.value = sqrt(arg.value); // calling double overload

    // reference record on tape
    result.idx = tape.size() - 1;

    // populate record on tape
    rec.numArg = 1;
    rec.idx1 = arg.idx;

    // compute derivative
    rec.der1 = 0.5 / result.value;

    return result;
}
```

# Custom function overloading

- Also overload building blocks that are not standard to C++ but frequently applied in applications
- In financial application, we use cumulative normal distributions and normal densities all the time
- The functions are defined in the file gaussians.h in the repo
- We overload here (once again, only change is in value and derivatives):

```
Number normalDens(const Number& arg)
{
    // create a new record on tape
    tape.push_back(Record());
    Record& rec = tape.back();

    // compute result
    Number result;
    result.value = normalDens(arg.value);

    // reference record on tape
    result.idx = tape.size() - 1;

    // populate record on tape
    rec.numArg = 1;
    rec.idx1 = arg.idx;

    // compute derivative
    rec.der1 = - result.value * arg.value;

    return result;
}
```

```
Number normalCdf(const Number& arg)
{
    // create a new record on tape
    tape.push_back(Record());
    Record& rec = tape.back();

    // compute result
    Number result;
    result.value = normalCdf(arg.value); // calling double overload in gaussians.h

    // reference record on tape
    result.idx = tape.size() - 1;

    // populate record on tape
    rec.numArg = 1;
    rec.idx1 = arg.idx;

    // compute derivative
    rec.der1 = normalDens(arg.value);

    return result;
}
```

# Avoiding code duplication

- The code for all binary operators and for all unary functions is identical
- Only values and derivatives are different
- This is poor design:  
If we change something in the logic, we must consistently modify many different functions
- In the professional code of chapters 10 and 15  
we structure code and apply “policy design” (Alexandresku, 2001) to avoid duplication without overhead
- Here, we stick with duplicated code

# Comparison operator overloading

- Our custom number must do everything a double does
- In particular, we must be able to compare two Numbers
- Hence, to complete our simple framework, we must also overload comparison operators:

```
bool operator==(const Number& lhs, const Number& rhs) { return lhs.value == rhs.value; }
bool operator!=(const Number& lhs, const Number& rhs) { return lhs.value != rhs.value; }
bool operator>(const Number& lhs, const Number& rhs) { return lhs.value > rhs.value; }
bool operator>=(const Number& lhs, const Number& rhs) { return lhs.value >= rhs.value; }
bool operator<(const Number& lhs, const Number& rhs) { return lhs.value < rhs.value; }
bool operator<=(const Number& lhs, const Number& rhs) { return lhs.value <= rhs.value; }
```

# Applying the recording framework

- Our simple recording framework is complete, see complete code in the GitHub repo, file toyCode.h
- We may use it to record calculations
- Example: Black and Scholes

```
inline double blackScholes(  
    // input layer 0  
    const double spot, const double rate, const double yield, const double vol, const double strike, const double mat)  
{  
    /* layer 1 */      double df = exp(-rate * mat), fwd = spot * exp((rate - yield) * mat), std = vol * sqrt(mat);  
    /* layer 2 */      double d = log(fwd / strike) / std;  
    /* layer 3 */      double d1 = d + 0.5 * std, d2 = d - 0.5 * std;  
    /* layer 4 */      double p1 = normalCdf(d1), p2 = normalCdf(d2);  
    /* output layer 5 */ return df * (fwd * p1 - strike * p2);  
}
```

# Instrumenting computation code

- To record a Black & Scholes calculation, we must call it with our number type
- This is called *instrumentation*
- We can replace all doubles by Numbers:

```
inline Number blackScholes(  
    // input layer 0  
    const Number spot, const Number rate, const Number yield, const Number vol, const Number strike, const Number mat)  
{  
    /* layer 1 */    Number df = exp(-rate * mat), fwd = spot * exp((rate - yield) * mat), std = vol * sqrt(mat);  
    /* layer 2 */    Number d = log(fwd / strike) / std;  
    /* layer 3 */    Number d1 = d + 0.5 * std, d2 = d - 0.5 * std;  
    /* layer 4 */    Number p1 = normalCdf(d1), p2 = normalCdf(d2);  
    /* output layer 5 */ return df * (fwd * p1 - strike * p2);  
}
```

# Instrumenting computation code

- Better solution: template code on number representation type

```
template <class T> inline T blackScholes(  
    // input layer 0  
    const T spot, const T rate, const T yield, const T vol, const T strike, const T mat)  
{  
    /* layer 1 */      T df = exp(-rate * mat), fwd = spot * exp((rate - yield) * mat), std = vol * sqrt(mat);  
    /* layer 2 */      T d = log(fwd / strike) / std;  
    /* layer 3 */      T d1 = d + 0.5 * std, d2 = d - 0.5 * std;  
    /* layer 4 */      T p1 = normalCdf(d1), p2 = normalCdf(d2);  
    /* output layer 5 */ return df * (fwd * p1 - strike * p2);  
}
```

- Best practice: produce templated code in the first place

- To evaluate only, call with doubles as arguments

```
double spot = 100, rate = 0.02, yield = 0.05, vol = 0.2, strike = 110, mat = 2; // initializes inputs  
auto result = blackScholes(spot, rate, yield, vol, strike, mat);                // evaluates operations
```

- To evaluate and record code, call with Numbers as arguments

```
Number spot = 100, rate = 0.02, yield = 0.05, vol = 0.2, strike = 110, mat = 2; // initializes and records inputs  
auto result = blackScholes(spot, rate, yield, vol, strike, mat);                // evaluates and records operations
```

AAD




# State of the tape after recording

- We call the instrumented instance of our code:

```
Number spot = 100, rate = 0.02, yield = 0.05, vol = 0.2, strike = 110, mat = 2; // initializes and records inputs
auto result = blackScholes(spot, rate, yield, vol, strike, mat); // evaluates and records operations
cout << result.value; // 5.03705
```

- Which evaluates the calculation and records all operations on tape
  - Denote  $f_i$  the operation number  $i$ , this is a function of 0, 1 or 2 arguments
  - The arguments must also be on tape, with indices  $< i$
  - Denote  $A_i$  the set of indices of the arguments to  $f_i$  - this is a set of 0, 1 or 2 indices, all  $< i$
  - Denote  $n$  the number of inputs and  $N$  the number of operations on tape, including inputs
  - Denote  $x_i$  the result of operation  $i$  and  $y = x_N$  the final result
- Note that all the local derivatives  $\frac{\partial f_i}{\partial x_j}$  for all  $j \in A_i$  have been computed and recorded on tape during evaluation


- Our tape therefore looks like:



# Successors and adjoints

- Successors

- Denote  $S_j$  the set of indices of the *successors* of  $x_j$  on tape
- These are the indices  $i$  of all functions  $f_i$ , calculated **after**  $x_j$  that use  $x_j$  as an argument
- Formally:  $S_j = \{i > j, j \in A_i\}$
- Note that  $S_i$  may contain many indices
- Whereas an empty  $S_i$  reveals an unused input or intermediate result



- Adjoint equation

- Denote  $\bar{x}_i \equiv \frac{\partial y}{\partial x_i} = \frac{\partial x_N}{\partial x_i}$  the adjoint of  $x_i$
- Then (evidently):  $\bar{x}_N = 1$
- And in a direct application of the chain rule:  $\bar{x}_j = \sum_{i \in S_j} \frac{\partial f_i}{\partial x_j} \bar{x}_i$  because  $\frac{\partial y}{\partial x_j} = \sum_{i \in S_j} \frac{\partial y}{\partial x_i} \frac{\partial x_i}{\partial x_j}$  and we recall that all the  $\frac{\partial x_i}{\partial x_j}$  are on tape

# Adjoint propagation

- Adjoints therefore satisfy a backward recursion
  - The last adjoint  $\bar{x}_N = 1$  is given
  - All other adjoints are a function of future adjoints: for all  $j$   $\bar{x}_j = f(\bar{x}_i, i \in S_j, i > j)$
- Therefore the following algorithm is guaranteed to correctly accumulate all adjoints
  - Initialize all adjoints to 0 and the last adjoint to 1 (this is called *seeding* the tape):  $\bar{x}_j = \delta_{N-j}$
  - Repeat for  $i$  iterating backwards from  $N$  to  $0$ : for all  $j \in A_i$  :  $\bar{x}_j \leftarrow \bar{x}_j + \frac{\partial f_i}{\partial x_j} \bar{x}_i$
  - The differentials of the calculation to its inputs  $x_1, x_2, \dots, x_n$  are, by definition,  $\bar{x}_1, \bar{x}_2, \dots, \bar{x}_n$
- This algorithm is called (reverse) adjoint propagation
- AAD is the sum of a recording framework and adjoint propagation

# Adjoint propagation code

```
vector<double> calculateAdjoints(Number& result)
{
    // initialization
    vector<double> adjoints(tape.size(), 0.0); // initialize all to 0
    int N = result.idx;                       // find N
    adjoints[N] = 1.0;                        // seed aN = 1

    // backward propagation
    for(int j=N; j>0; --j) // iterate backwards over tape
    {
        if (tape[j].numArg > 0)
        {
            adjoints[tape[j].idx1] += adjoints[j] * tape[j].der1; // propagate first argument

            if (tape[j].numArg > 1)
            {
                adjoints[tape[j].idx2] += adjoints[j] * tape[j].der2; // propagate second argument
            }
        }
    }

    return adjoints;
}
```

# Complexity

- One *evaluation* of the calculation
  - Sweeps forward through the sequence of its operations
  - Executes every operation exactly once
  - Therefore its complexity is  $N$ , the number of records that end up on tape
- Adjoint propagation
  - Also sweeps through the sequence of operations, backward
  - Executes 0, 1 or 2 operations on every record, depending on the number of arguments
  - Therefore, as advertised, AAD computes **all  $n$**  differentials in **constant time**
  - In theory, all differentials are propagated in less than 2x one evaluation
  - In addition, the calculation must be evaluated (and recorded) first so the theoretical upper bound is 3x one evaluation
  - Due to recording and tape traversal overhead, a good implementation generally produces many differentials in 4x to 10x
- Our professional code from chapters 10, 12 and 15 beats the theoretical bound!
  - Recall from the demonstration, one evaluation = 0.8sec, 1,081 differentials = 1.5sec, less than 2x one evaluation
  - Due to “selective instrumentation”, a strong optimization, explained, along with many others, in chapter 12

# Adjoint propagation code

- After we record a tape by calling the instrumented instance of our code:

```
Number spot = 100, rate = 0.02, yield = 0.05, vol = 0.2, strike = 110, mat = 2; // initializes and records inputs
auto result = blackScholes(spot, rate, yield, vol, strike, mat); // evaluates and records operations
cout << result.value; // 5.03705
```

- We proceed with adjoint propagation:

```
// propagate adjoints
vector<double> adjoints = calculateAdjoints(result);
```

- The code works nicely  
and produces correct values

```
// show derivatives
cout << "Derivative to spot (delta) = " << adjoints[spot.idx] << endl; // 0.309
cout << "Derivative to rate (rho) = " << adjoints[rate.idx] << endl; // 51.772
cout << "Derivative to dividend yield = " << adjoints[yield.idx] << endl; // -61.846
cout << "Derivative to volatility (vega) = " << adjoints[vol.idx] << endl; // 46.980
cout << "Derivative to strike (-digital) = " << adjoints[strike.idx] << endl; // -0.235
cout << "Derivative to maturity (-theta) = " << adjoints[mat.idx] << endl; // 1.321
```

- Not very interesting for Black & Scholes
  - Fast, analytic evaluation
  - Only 6 differentials to compute
- Next, we apply AAD to a barrier option in a Monte-Carlo simulation of Dupire's model

# Conclusion

- We implemented AAD in the simplest possible manner, scratching the surface of possibilities
- Part III (Chapters 8 to 15) gives the details of a complete, professional, efficient implementation
  - How to minimize recording overhead
  - Efficient memory management constructs
  - Apply check-pointed AAD to differentiate a calculation piece by piece to mitigate RAM footprint and cache inefficiency
  - Efficiently differentiate non-scalar calculations that return multiple results
  - Cutting-edge implementation with template meta-programming and expression templates, faster by 2x to 5x
  - Parallel implementation
  - Advice for debugging and optimization
  - And much more
- Still the simplistic code works and produces the correct values
- Not very interesting for Black & Scholes
  - Fast, analytic evaluation
  - Only 6 differentials to compute
- Next, we apply the framework to a barrier option in Dupire Monte-Carlo
  - Long, complex evaluation
  - 1,081 differentials to compute

AAD for financial simulations



# Simple simulation code

- We implement a simplistic simulation code for a barrier option in Dupire's model
  - Recall local volatility is given in a matrix and bi-linearly interpolated in spot and time
- We need the following pieces, which we assume are given here
  - A matrix class to hold local volatilities (matrix.h in the repo, chapters 1 and 2 in the book)
  - A bi-linear interpolation function (interp.h in the repo, chapter 6, section 6.4 in the book)
  - Random number generators to produce independent Gaussian increments (chapters 5 and 6)
- The code is templated on the real number representation type

# Simulation code, version 1

// Signature

```
template <class T>
inline T toyDupireBarrierMc(
    // Spot
    const T S0,
    // Local volatility
    const vector<T> spots,
    const vector<T> times,
    const matrix<T> vols,
    // Product parameters
    const T maturity,
    const T strike,
    const T barrier,
    // Number of paths and time steps
    const int Np,
    const int Nt,
    // Initialized random number generator
    RNG& random)
```

// Implementation

```
// Initialize
T result = 0;
// double because the RNG is not templated (and doesn't need to be, see chapter 12)
vector<double> gaussianIncrements(Nt);
const T dt = maturity / Nt, sdt = sqrt(dt);

// Loop over paths
for (int i = 0; i < Np; ++i)
{
    // Generate Nt Gaussian Numbers
    random.nextG(gaussianIncrements);
    // Euler's scheme, step by step
    T spot = S0, time = 0;
    bool alive = true;
    for (size_t j = 0; j < Nt; ++j)
    {
        // Interpolate volatility
        const T vol = interp2D(spots, times, vols, spot, time);
        time += dt;
        // Simulate return
        spot *= exp(-0.5 * vol * vol * dt + vol * sdt * gaussianIncrements[j]);
        // Monitor barrier
        if (spot > barrier)
        {
            alive = false;
            break;
        }
    }
    // Payoff
    if (alive && spot > strike) result += spot - strike;
} // paths

return result / Np;
```

# A simplistic code

- This code “does the job” but is not acceptable by professional standards
- The code is specific to Dupire’s model and an up & out call, therefore not scalable
  - To price another product (Asian option, Ratchet option, ...) copy the code and change the lines that evaluate payoffs
  - To price in another model (Heston, SLV, ...) copy the code and change the lines that generate the scenarios
  - End up with many different functions implementing the same simulation logic for different couples of models and products
  - To modify the simulation logic, consistently change all the functions! This is obviously not viable
- Chapter 6 teaches a professional architecture for generic simulation libraries
  - Encapsulate scenario generation in Model objects
  - Encapsulate payoff evaluation in Product objects
  - Encapsulate simulation logic in a generic Monte-Carlo engine
  - Code every model and every product exactly once, mix and match at run time
- We stick with the simplistic code for demonstration purposes

# An inefficient code

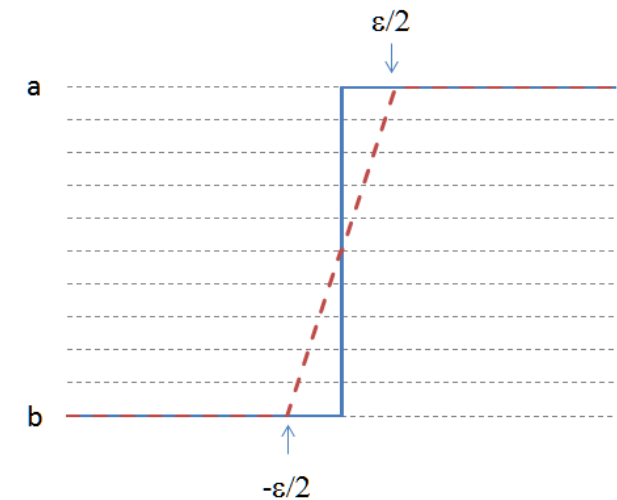
- The code does too much work repeatedly during simulations
  - Key to efficient Monte-Carlo code:  
**Do as much work as possible once, on initialization, and as little as possible repeatedly, during simulations**
  - Example: we perform an expensive bi-linear interpolation in the innermost loop, for every path, on every time step
  - We interpolate in spot and time, spot is stochastic (scenario dependent), time is not
  - Therefore we can (and should) pre-interpolate in time on initialization  
And perform only 1D interpolations in the innermost loop
- Chapter 6 teaches and builds fully optimized code
- The code is serial, it executes sequentially on one core
  - Since even phones are multi-core today, professional code is always parallel
  - With Monte-Carlo simulations, it is relatively easy to obtain a speed-up by the number of physical cores
- Chapter 3 teaches modern parallel C++, chapter 7 builds a professional *parallel* simulation library  
And section 12.5 instruments it with AAD in parallel
- For the purpose of demonstration, we stick with the not so efficient, serial version

# Smoothing barrier options

- Discretely monitored barrier options are discontinuous, their value jumps to 0 at the barrier
  - Therefore, our code is not differentiable
  - AAD does not help: it cannot perform the impossible task of differentiating a discontinuous functions
  - With finite differences, barrier risks are unstable, with AAD they are all zero
  - Find the reason why as an exercise!
- Therefore traders always smooth discontinuous transactions (barriers, digitals etc.)
- Smoothing = applying a close, continuous approximation in place of the discontinuous function
- Smoothing in finance, and its connection to fuzzy logic, are explained in the presentation:  
<http://www.slideshare.net/AntoineSavine/stabilise-risks-of-discontinuous-payoffs-with-fuzzy-logic>  
freely available on slideShare
- Here, we briefly explain the “smooth barrier” algorithm, universally applied on derivatives desks

# Smooth barrier

- Hard barrier
  - 100% dead above the barrier, 100% alive below the barrier
  - Hence, discontinuous
- Soft barrier
  - 100% dead above the barrier **plus epsilon**, 100% alive below the barrier **minus epsilon**
  - In between, lose a fraction of notional interpolated between (barrier-epsilon,0) and (barrier+epsilon,1)
  - And continue with the remaining notional
  - Hence, continuous
- Smoothing and fuzzy logic
  - Like Schrodinger's cat, the transaction is in a superposition of dead and alive states
  - Smoothing is achieved by replacing sharp logic (dead or alive?) by fuzzy logic (how much alive?)
  - More in the presentation



# Simulation code with smooth barrier

```
template <class T>
inline T toyDupireBarrierMc(
    // Spot
    const T S0,
    // Local volatility
    const vector<T> spots,
    const vector<T> times,
    const matrix<T> vols,
    // Product parameters
    const T maturity,
    const T strike,
    const T barrier,
    // Number of paths and time steps
    const int Np,
    const int Nt,
    // Smoothing
    const T epsilon,
    // initialized random number generator
    RNG& random)

// Initialize
T result = 0;
// double because the RNG is not templated (and doesn't need to be, see chapter 12)
vector<double> gaussianIncrements(Nt);
const T dt = maturity / Nt, sdt = sqrt(dt);

// Loop over paths
for (int i = 0; i < Np; ++i)
{
    // Generate Nt Gaussian Numbers
    random.nextG(gaussianIncrements);
    // Step by step
    T spot = S0, time = 0;
    /* bool alive = true; */ T alive = 1.0; // alive is a real number in (0,1)
    for (size_t j = 0; j < Nt; ++j)
    {
        // Interpolate volatility
        const T vol = interp2D(spots, times, vols, spot, time);
        time += dt;
        // Simulate return
        spot *= exp(-0.5 * vol * vol * dt + vol * sdt * gaussianIncrements[j]);
        // Monitor barrier
        /* if (spot > barrier) { alive = false; break; } */
        if (spot > barrier + epsilon) { alive = 0.0; break; } // definitely dead
        else if (spot < barrier - epsilon) { /* do nothing */ }; // definitely alive
        else /* in between, interpolate */ alive *= 1.0 - (spot - barrier + epsilon) / (2 * epsilon);
    }
    // Payoff paid on surviving notional
    /* if (alive && spot > strike) result += spot - strike; */ if (spot > strike) result += alive * (spot - strike);
} // paths
return result / Np;
```

# Simulation code

- The toy simulation code is found on the file ToyCode.h in the repo
- To be compared with professional code in the files with names prefixed by “mc”
- Our simple code returns the same result as the professional code
- It is twice slower than the serial version of the professional code
- On a quad-core computer, it is 8x slower than the parallel version
- Next, we differentiate it with our simple AAD framework



# Differentiation

Just like we did for Black & Scholes, to compute differentials, we:

1. Initialize the inputs as Numbers  
Which also records them on tape
2. Call our templated evaluation code, instantiated with the Number type  
Which performs the evaluation *and* records operations on tape
3. Propagate adjoints backwards through the tape
4. Pick differentials as the adjoints of the parameters

# Differentiation code

```
void toyDupireBarrierMcRisks(
    const double S0, const vector<double> spots, const vector<double> times, const matrix<double> vols,
    const double maturity, const double strike, const double barrier,
    const int Np, const int Nt, const double epsilon, RNG& random,
    /* results: value and dV/dS, dV/d(local vols) */ double& price, double& delta, matrix<double>& vegas)
{
    // 1. Initialize inputs

    Number nS0(S0), nMaturity(maturity), nStrike(strike), nBarrier(barrier), nEpsilon(epsilon);
    vector<Number> nSpots(spots.size()), nTimes(times.size());
    matrix<Number> nVols(vols.rows(), vols.cols());

    for (int i = 0; i < spots.size(); ++i) nSpots[i] = Number(spots[i]);
    for (int i = 0; i < times.size(); ++i) nTimes[i] = Number(times[i]);
    for (int i = 0; i < vols.rows(); ++i) for (int j = 0; j < vols.cols(); ++j) nVols[i][j] = Number(vols[i][j]);

    // 2. Call instrumented evaluation code, which evaluates the barrier option price and records all operations
    Number nPrice = toyDupireBarrierMc(nS0, nSpots, nTimes, nVols, nMaturity, nStrike, nBarrier, Np, Nt, nEpsilon, random);

    // 3. Adjoint propagation, the exact same code as before, should be encapsulated in a dedicated function
    vector<double> adjoints = calculateAdjoints(nPrice);

    // 4. Pick results
    price = nPrice.value;
    delta = adjoints[nS0.idx];
    for (int i = 0; i < vols.rows(); ++i) for (int j = 0; j < vols.cols(); ++j) vegas[i][j] = adjoints[nVols[i][j].idx];
}
```

# Testing the code

- We run the code in the same context as the initial demonstration but with 100,000 paths instead of 500,000
- The computer runs out of memory and crashes!
- Running AAD on a simulation with 100,000 paths consumes an insane amount of RAM
- Even on a computer with enough memory, such large tape is cache inefficient

# Solution in principle

- Run a series of risks on mini-batches of say, 1024 paths and average in the end
- Wipe the tape in between mini-batches
- The average of differentials is the differential of the average
- So we get the same results while reducing memory footprint to operations recorded over 1,024 paths
- Note with mini-batches of size 1, this is known as “path-wise differentiation”
- This is also a particular, and simple case of the general check-pointing algorithm explained in chapter 13

# Solution in code

- Rename our function DupireRisksMiniBatch(), call it sequentially from a wrapper function

```
void toyDupireBarrierMcRisks(
    const double S0, const vector<double> spots, const vector<double> times, const matrix<double> vols,
    const double maturity, const double strike, const double barrier,
    const int Np, const int Nt, const double epsilon, RNG& random,
    /* results: value and dV/dS, dV/d(local vols) */ double& price, double& delta, matrix<double>& vegas)
{
    price = delta = 0;
    for (int i = 0; i < vegas.rows(); ++i) for (int j = 0; j < vegas.cols(); ++j) vegas[i][j] = 0;
    double batchPrice, batchDelta; matrix<double> batchVegas(vegas.rows(), vegas.cols());
    int pathsToGo = Np, pathsPerBatch = 1024;
    // calculate batch sensitivities sequentially
    while (pathsToGo > 0)
    {
        // wipe tape
        tape.clear();

        // do mini batch
        int paths = min(pathsToGo, pathsPerBatch);
        dupireRisksMiniBatch(S0, spots, times, vols, maturity, strike, barrier, paths, Nt, epsilon, random, batchPrice, batchDelta, batchVegas);

        // update results
        price += batchPrice * paths / Np;
        delta += batchDelta * paths / Np;
        for (int i = 0; i < vegas.rows(); ++i) for (int j = 0; j < vegas.cols(); ++j) vegas[i][j] += batchVegas[i][j] * paths / Np;


        pathsToGo -= paths;
    }
}
```

# Performance

- With 100,000 paths, 156 steps, we compute the 1,081 differentials in around 7 seconds
- This is 1,081 differentials in the time of around 6 evaluations
- This is a very remarkable result, especially with such simplistic code
- Try it yourself with the code in the repo!
- This being said, the professional code is around 8 times faster in serial mode, 32 times faster in parallel mode on a quad-core laptop

# Conclusion

- We learned AAD in principle and in code  
And applied it to machine learning and finance
- But we really just scratched the surface
- For example:
  - How to efficiently differentiate the multiple results of non-scalar functions?
  - How to compute risks not on model parameters like local volatilities  
But on tradable market variables like implied volatilities?
  - How to implement AAD in modern C++ and manage memory efficiently?
  - How to implement AAD over parallel simulations and run it at least 32x faster?
- The answers, and much more, are in the book



Thank you for your attention. Find the slides here:

<http://github.com/asavine/CompFinance/Intro2AADinMachineLearningAndFinance.pdf>

<http://tinyurl.com/savineAadTalk>

follow asavine on GitHub and watch the repo to be notified of updates